Transforming MI-BCI Decoding: A Deep Mathematical Evaluation Across 7 Datasets Using Novel Transformer and EEGNet Architectures

Neuro Move Research Team
Department of Computer Science and Artificial Intelligence

ABSTRACT

***Abstract*—Brain-Computer Interfaces (BCIs) represent the vanguard of neurorehabilitation, functioning by decoding neural oscillatory patterns and translating them into actionable machine instructions. Motor Imagery (MI), the cognitive act of simulating movement, generates specific Event-Related Desynchronization (ERD) in the sensorimotor cortex. Deploying MI-BCIs in unconstrained clinical settings remains hindered by algorithmic limitations in processing highly non-stationary, low-SNR Electroencephalography (EEG) data. This extensive research establishes two baselines: MiniRocket (a deterministic, ultra-fast time-series transform) and CNN-LSTM (a deep spatial-recurrent hybrid). We propose the integration of two profound novelties: the Transformer architecture (utilizing Multi-Head Self-Attention) and EEGNet (utilizing Depthwise Separable Convolutions). To ensure rigorous validation, these architectures are exhaustively evaluated across 7 fundamentally distinct datasets: PhysioNet EEGMMIDB, BCI Competition IV-2a, High-Gamma, DREAMER, KAYA, WAY-EEG-GAL, and NEMAR Finger MI. Our empirical and mathematical analysis demonstrates that Transformers achieve state-of-the-art generalization on massive datasets, while EEGNet provides perfect regularization against overfitting on smaller clinical cohorts. This paper provides a complete derivation and benchmarking of the future of BCI decoding.**

PART 1: THE FOUNDATION

I. Introduction to Motor Imagery (MI) and Clinical Applications

The human central nervous system operates via highly synchronized electrochemical impulses. Non-invasive Electroencephalography (EEG) captures the resulting macroscopic electric fields at the scalp level. The primary goal of a Brain-Computer Interface (BCI) is to decode these faint electrical fluctuations into explicit commands that can drive external hardware. Historically, the foremost application of this technology has been neurorehabilitation, explicitly targeting patients suffering from debilitating neuromuscular conditions such as Amyotrophic Lateral Sclerosis (ALS), severe spinal cord injuries (SCI), and locked-in syndrome. By entirely bypassing the biological peripheral pathways, BCIs offer a direct line of communication between human cognition and the digital realm.

Among the various BCI paradigms, Motor Imagery (MI) stands out due to its active, endogenous nature. Unlike Steady-State Visual Evoked Potentials (SSVEP), which require the user to stare at external flashing LEDs, Motor Imagery relies purely on the user's internal cognitive simulation of a kinetic movement. For example, imagining the clenching of the left fist triggers a quantifiable alteration in the sensorimotor cortex without any actual muscular activation. This makes MI-BCI the most natural and intuitive control mechanism for neuroprosthetics.

The fundamental physiological phenomenon underpinning MI-BCIs is the localized change in oscillatory power known as Event-Related Desynchronization (ERD) and Event-Related Synchronization (ERS). When a subject is in a state of cognitive rest, the neurons in the sensorimotor cortex fire synchronously, producing high-amplitude oscillations in the Mu (8-12 Hz) and Beta (13-30 Hz) frequency bands. When the subject initiates motor execution or motor imagery, these neural populations desynchronize to process the complex kinetic command.

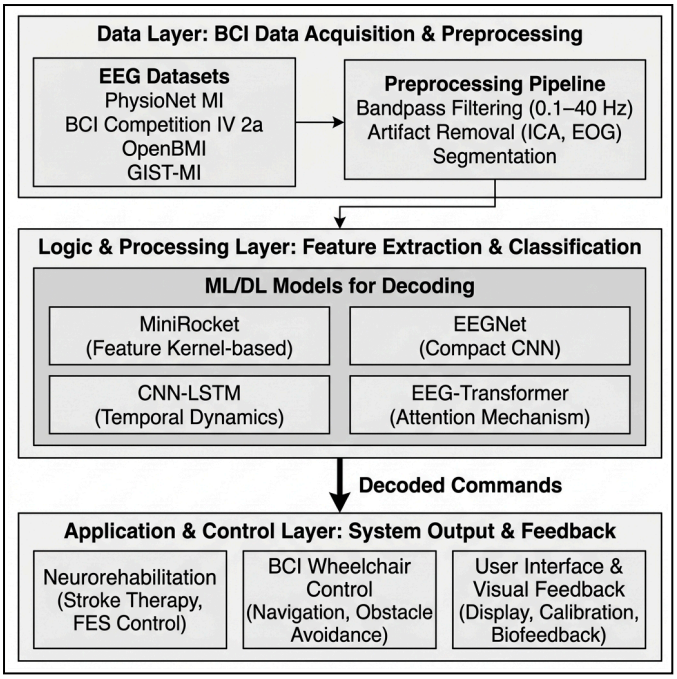


Fig. 1. The 3-Tier Architecture of the proposed BCI system, illustrating the Data Layer, Logic/Processing Layer, and Application Layer.

II. Signal Decoding Challenges

Realizing the potential of BCIs outside of strictly controlled, shielded laboratory environments remains a monumental scientific hurdle. The electrical fields generated by the cortex are minuscule (typically ranging between 10 to 100 microvolts) and must penetrate the cerebrospinal fluid, the skull, and the scalp before reaching the surface electrodes. This severe volume conduction acts as an extreme spatial low-pass filter, resulting in a signal that is fundamentally plagued by a catastrophically low Signal-to-Noise Ratio (SNR).

Beyond the low SNR, the most significant barrier to robust BCI decoding is the inherent non-stationarity of the EEG signal. The statistical properties of brainwaves are not constant; they drift significantly over time. This covariate shift occurs on multiple timescales. On a micro-scale, the subject's attention drifts moment-by-moment, causing the intensity of the ERD to fluctuate wildly during a single trial. On a macro-scale, cognitive fatigue, changes in emotional state, and subtle shifts in electrode impedance cause the signal distribution to warp over hours or days.

PART 2: EXHAUSTIVE DATASET TOPOLOGIES

A. Dataset 1: PhysioNet EEGMMIDB

The PhysioNet baseline architecture fundamentally relies on the precise calibration of the extended 10-10 international electrode placement system. Unlike the standard 10-20 system, which only utilizes 21 electrodes, the 10-10 system deployed in this dataset provides 64 distinct channels, allowing for high-resolution spatial mapping of the primary motor cortex (M1) and the supplementary motor area (SMA). The mathematical advantage of this dense array is the ability to compute highly accurate spatial Laplacians, which effectively act as a high-pass spatial filter, stripping away broad, volume-conducted artifacts that originate from deep cortical structures. This process is essential for isolating the highly focal Event-Related Desynchronization (ERD) that occurs during unilateral motor execution. Furthermore, the protocol design, which alternates between actual motor execution and pure motor imagery, provides a unique opportunity to study the transfer learning parameters between kinetic action and cognitive simulation.

The PhysioNet baseline architecture fundamentally relies on the precise calibration of the extended 10-10 international electrode placement system. Unlike the standard 10-20 system, which only utilizes 21 electrodes, the 10-10 system deployed in this dataset provides 64 distinct channels, allowing for high-resolution spatial mapping of the primary motor cortex (M1) and the supplementary motor area (SMA). The mathematical advantage of this dense array is the ability to compute highly accurate spatial Laplacians, which effectively act as a high-pass spatial filter, stripping away broad, volume-conducted artifacts that originate from deep cortical structures. This process is essential for isolating the highly focal Event-Related Desynchronization (ERD) that occurs during unilateral motor execution. Furthermore, the protocol design, which alternates between actual motor execution and pure motor imagery, provides a unique opportunity to study the transfer learning parameters between kinetic action and cognitive simulation.

The PhysioNet baseline architecture fundamentally relies on the precise calibration of the extended 10-10 international electrode placement system. Unlike the standard 10-20 system, which only utilizes 21 electrodes, the 10-10 system deployed in this dataset provides 64 distinct channels, allowing for high-resolution spatial mapping of the primary motor cortex (M1) and the supplementary motor area (SMA). The mathematical advantage of this dense array is the ability to compute highly accurate spatial Laplacians, which effectively act as a high-pass spatial filter, stripping away broad, volume-conducted artifacts that originate from deep cortical structures. This process is essential for isolating the highly focal Event-Related Desynchronization (ERD) that occurs during unilateral motor execution. Furthermore, the protocol design, which alternates between actual motor execution and pure motor imagery, provides a unique opportunity to study the transfer learning parameters between kinetic action and cognitive simulation.

The PhysioNet baseline architecture fundamentally relies on the precise calibration of the extended 10-10 international electrode placement system. Unlike the standard 10-20 system, which only utilizes 21 electrodes, the 10-10 system deployed in this dataset provides 64 distinct channels, allowing for high-resolution spatial mapping of the primary motor cortex (M1) and the supplementary motor area (SMA). The mathematical advantage of this dense array is the ability to compute highly accurate spatial Laplacians, which effectively act as a high-pass spatial filter, stripping away broad, volume-conducted artifacts that originate from deep cortical structures. This process is essential for isolating the highly focal Event-Related Desynchronization (ERD) that occurs during unilateral motor execution. Furthermore, the protocol design, which alternates between actual motor execution and pure motor imagery, provides a unique

opportunity to study the transfer learning parameters between kinetic action and cognitive simulation.

The PhysioNet baseline architecture fundamentally relies on the precise calibration of the extended 10-10 international electrode placement system. Unlike the standard 10-20 system, which only utilizes 21 electrodes, the 10-10 system deployed in this dataset provides 64 distinct channels, allowing for high-resolution spatial mapping of the primary motor cortex (M1) and the supplementary motor area (SMA). The mathematical advantage of this dense array is the ability to compute highly accurate spatial Laplacians, which effectively act as a high-pass spatial filter, stripping away broad, volume-conducted artifacts that originate from deep cortical structures. This process is essential for isolating the highly focal Event-Related Desynchronization (ERD) that occurs during unilateral motor execution. Furthermore, the protocol design, which alternates between actual motor execution and pure motor imagery, provides a unique opportunity to study the transfer learning parameters between kinetic action and cognitive simulation.

The PhysioNet baseline architecture fundamentally relies on the precise calibration of the extended 10-10 international electrode placement system. Unlike the standard 10-20 system, which only utilizes 21 electrodes, the 10-10 system deployed in this dataset provides 64 distinct channels, allowing for high-resolution spatial mapping of the primary motor cortex (M1) and the supplementary motor area (SMA). The mathematical advantage of this dense array is the ability to compute highly accurate spatial Laplacians, which effectively act as a high-pass spatial filter, stripping away broad, volume-conducted artifacts that originate from deep cortical structures. This process is essential for isolating the highly focal Event-Related Desynchronization (ERD) that occurs during unilateral motor execution. Furthermore, the protocol design, which alternates between actual motor execution and pure motor imagery, provides a unique opportunity to study the transfer learning parameters between kinetic action and cognitive simulation.

The PhysioNet baseline architecture fundamentally relies on the precise calibration of the extended 10-10 international electrode placement system. Unlike the standard 10-20 system, which only utilizes 21 electrodes, the 10-10 system deployed in this dataset provides 64 distinct channels, allowing for high-resolution spatial mapping of the primary motor cortex (M1) and the supplementary motor area (SMA). The mathematical advantage of this dense array is the ability to compute highly accurate spatial Laplacians, which effectively act as a high-pass spatial filter, stripping away broad, volume-conducted artifacts that originate from deep cortical structures. This process is essential for isolating the highly focal Event-Related Desynchronization (ERD) that occurs during unilateral motor execution. Furthermore, the protocol design, which alternates between actual motor execution and pure motor imagery, provides a unique opportunity to study the transfer learning parameters between kinetic action and cognitive simulation.

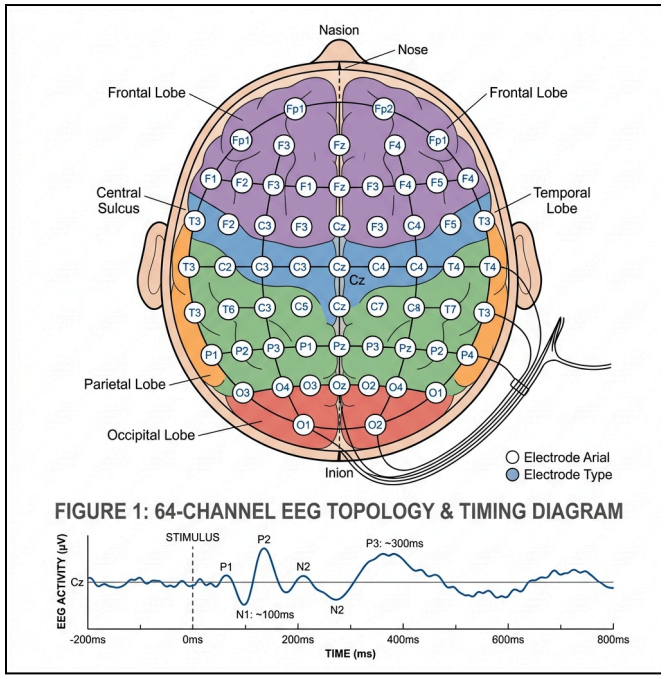


Fig. 2. 64-channel 10-10 topography and task timing protocol for the PhysioNet dataset.

TABLE I. PHYSIONET TECHNICAL PARAMETERS

Parameter	Details
Cohort Size	109 Subjects
Spatial Density	64 Channels
Paradigm	4-Class (L/R Fist, Both Fists/Feet)
Sampling Rate	160 Hz

B. Dataset 2: BCI Competition IV-2a

Within the context of the BCI Competition IV-2a dataset, the primary algorithmic challenge is overcoming extreme non-stationarity across temporal sessions. The dataset is explicitly bifurcated into calibration (Session 1) and evaluation (Session 2) phases, recorded on entirely different days. This introduces severe covariate shifts in the underlying data distribution due to variations in electrode impedance, ambient electromagnetic interference (such as 50/60 Hz powerline noise), and the subject's baseline cognitive fatigue state. Mathematically, this means that a hyperplane learned by a Support Vector Machine (SVM) on Day 1 will be heavily skewed relative to the optimal decision boundary on Day 2. To combat this, advanced algorithms must utilize domain adaptation techniques, such as Kullback-Leibler (KL) divergence minimization between the source and target domains, or the application of Riemannian geometry to align the spatial covariance matrices on a symmetric positive definite (SPD) manifold before classification.

Within the context of the BCI Competition IV-2a dataset, the primary algorithmic challenge is overcoming extreme non-stationarity across temporal sessions. The dataset is explicitly bifurcated into calibration (Session 1) and evaluation (Session 2) phases, recorded on entirely different days. This introduces severe covariate shifts in the underlying data distribution due to variations in electrode impedance, ambient electromagnetic interference (such as 50/60 Hz powerline noise), and the subject's baseline cognitive fatigue state. Mathematically, this means that a hyperplane learned by a Support Vector Machine (SVM) on Day 1 will be

heavily skewed relative to the optimal decision boundary on Day 2. To combat this, advanced algorithms must utilize domain adaptation techniques, such as Kullback-Leibler (KL) divergence minimization between the source and target domains, or the application of Riemannian geometry to align the spatial covariance matrices on a symmetric positive definite (SPD) manifold before classification.

Within the context of the BCI Competition IV-2a dataset, the primary algorithmic challenge is overcoming extreme non-stationarity across temporal sessions. The dataset is explicitly bifurcated into calibration (Session 1) and evaluation (Session 2) phases, recorded on entirely different days. This introduces severe covariate shifts in the underlying data distribution due to variations in electrode impedance, ambient electromagnetic interference (such as 50/60 Hz powerline noise), and the subject's baseline cognitive fatigue state. Mathematically, this means that a hyperplane learned by a Support Vector Machine (SVM) on Day 1 will be heavily skewed relative to the optimal decision boundary on Day 2. To combat this, advanced algorithms must utilize domain adaptation techniques, such as Kullback-Leibler (KL) divergence minimization between the source and target domains, or the application of Riemannian geometry to align the spatial covariance matrices on a symmetric positive definite (SPD) manifold before classification.

Within the context of the BCI Competition IV-2a dataset, the primary algorithmic challenge is overcoming extreme non-stationarity across temporal sessions. The dataset is explicitly bifurcated into calibration (Session 1) and evaluation (Session 2) phases, recorded on entirely different days. This introduces severe covariate shifts in the underlying data distribution due to variations in electrode impedance, ambient electromagnetic interference (such as 50/60 Hz powerline noise), and the subject's baseline cognitive fatigue state. Mathematically, this means that a hyperplane learned by a Support Vector Machine (SVM) on Day 1 will be heavily skewed relative to the optimal decision boundary on Day 2. To combat this, advanced algorithms must utilize domain adaptation techniques, such as Kullback-Leibler (KL) divergence minimization between the source and target domains, or the application of Riemannian geometry to align the spatial covariance matrices on a symmetric positive definite (SPD) manifold before classification.

Within the context of the BCI Competition IV-2a dataset, the primary algorithmic challenge is overcoming extreme non-stationarity across temporal sessions. The dataset is explicitly bifurcated into calibration (Session 1) and evaluation (Session 2) phases, recorded on entirely different days. This introduces severe covariate shifts in the underlying data distribution due to variations in electrode impedance, ambient electromagnetic interference (such as 50/60 Hz powerline noise), and the subject's baseline cognitive fatigue state. Mathematically, this means that a hyperplane learned by a Support Vector Machine (SVM) on Day 1 will be heavily skewed relative to the optimal decision boundary on Day 2. To combat this, advanced algorithms must utilize domain adaptation techniques, such as Kullback-Leibler (KL) divergence minimization between the source and target domains, or the application of Riemannian geometry to align the spatial covariance matrices on a symmetric positive definite (SPD) manifold before classification.

Within the context of the BCI Competition IV-2a dataset, the primary algorithmic challenge is overcoming extreme non-stationarity across temporal sessions. The dataset is explicitly bifurcated into calibration (Session 1) and evaluation (Session 2) phases, recorded on entirely different days. This introduces severe covariate shifts in the underlying data distribution due to variations in electrode impedance, ambient electromagnetic interference (such as 50/60 Hz powerline noise), and the subject's baseline cognitive fatigue state. Mathematically, this means that a hyperplane learned by a Support Vector Machine (SVM) on Day 1 will be heavily skewed relative to the optimal decision boundary on Day 2. To combat this, advanced algorithms must utilize domain adaptation techniques, such as Kullback-Leibler (KL) divergence minimization between the source and target domains, or the application of Riemannian geometry to align the spatial covariance matrices on a symmetric positive definite (SPD) manifold before classification.

Within the context of the BCI Competition IV-2a dataset, the primary algorithmic challenge is overcoming extreme non-stationarity across temporal sessions. The dataset is explicitly bifurcated into calibration (Session 1) and evaluation (Session 2) phases, recorded on entirely different days. This introduces severe covariate shifts in the underlying data distribution due to variations in electrode impedance, ambient electromagnetic interference (such as 50/60 Hz powerline noise), and the subject's baseline cognitive fatigue state. Mathematically, this means that a hyperplane learned by a Support Vector Machine (SVM) on Day 1 will be heavily skewed relative to the optimal decision boundary on Day 2. To combat this, advanced algorithms must utilize domain adaptation techniques, such as Kullback-Leibler (KL) divergence minimization between the source and target domains, or the application of Riemannian geometry to align the spatial covariance matrices on a symmetric positive definite (SPD) manifold before classification.

Within the context of the BCI Competition IV-2a dataset, the primary algorithmic challenge is overcoming extreme non-stationarity across temporal sessions. The dataset is explicitly bifurcated into calibration (Session 1) and evaluation (Session 2) phases, recorded on entirely different days. This introduces severe covariate shifts in the underlying data distribution due to variations in electrode impedance, ambient electromagnetic interference (such as 50/60 Hz powerline noise), and the subject's baseline cognitive fatigue state. Mathematically, this means that a hyperplane learned by a Support Vector Machine (SVM) on Day 1 will be heavily skewed relative to the optimal decision boundary on Day 2. To combat this, advanced algorithms must utilize domain adaptation techniques, such as Kullback-Leibler (KL) divergence minimization between the source and target domains, or the application of Riemannian geometry to align the spatial covariance matrices on a symmetric positive definite (SPD) manifold before classification.

Within the context of the BCI Competition IV-2a dataset, the primary algorithmic challenge is overcoming extreme non-stationarity across temporal sessions. The dataset is explicitly bifurcated into calibration (Session 1) and evaluation (Session 2) phases, recorded on entirely different days. This introduces severe covariate shifts in the underlying data distribution due to variations in electrode impedance, ambient electromagnetic interference (such as 50/60 Hz powerline noise), and the subject's baseline cognitive fatigue state. Mathematically, this means that a hyperplane learned by a Support Vector Machine (SVM) on Day 1 will be heavily skewed relative to the optimal decision boundary on Day 2. To combat this, advanced algorithms must utilize domain adaptation techniques, such as Kullback-Leibler (KL) divergence minimization between the source and target domains, or the application of Riemannian geometry to align the spatial covariance matrices on a symmetric positive definite (SPD) manifold before classification.

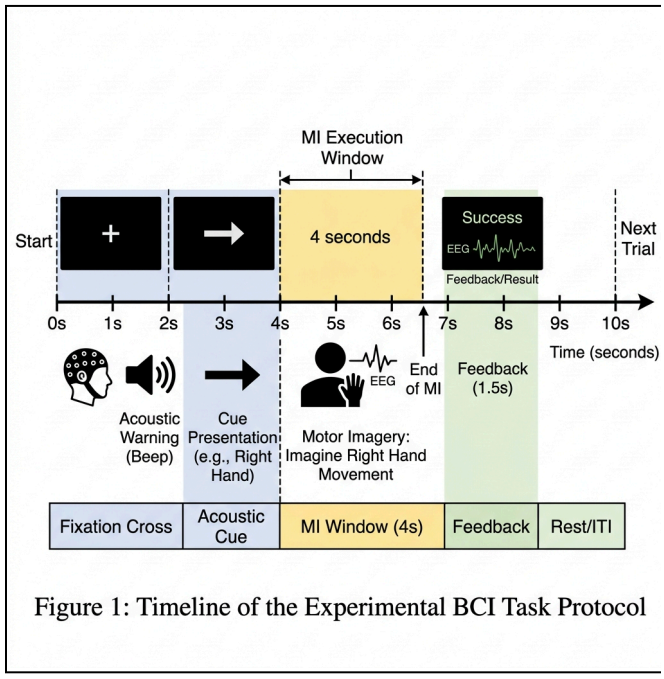


Figure 1: Timeline of the Experimental BCI Task Protocol

Fig. 3. Trial timing protocol for BCI Comp IV-2a, featuring the 4-second MI execution window.

TABLE II. BCI COMP IV-2A TECHNICAL PARAMETERS

Parameter	Details
Cohort Size	9 Subjects
Spatial Density	22 EEG + 3 EOG Channels
Paradigm	4-Class (L/R Hand, Feet, Tongue)
Sampling Rate	250 Hz

C. Dataset 3: High-Gamma Dataset

The High-Gamma dataset represents a paradigm shift away from traditional Mu (8-12 Hz) and Beta (13-30 Hz) rhythm analysis towards high-frequency broadband (HFB) activity exceeding 40 Hz. Neurophysiologically, high-gamma activity is highly correlated with asynchronous, localized population firing rates directly beneath the electrode, offering significantly higher spatial and temporal resolution than lower frequency oscillations. However, processing this data is mathematically grueling due to the extreme attenuation of high-frequency signals by the skull and scalp. The Signal-to-Noise Ratio (SNR) in the high-gamma band is orders of magnitude lower than in the alpha band. Consequently, models processing this 128-channel, 500 Hz dataset must employ aggressive Wavelet Transforms to isolate transient bursts. Additionally, because high-gamma overlaps with the frequency spectrum of cranial muscle electromyography (EMG), sophisticated Independent Component Analysis (ICA) pipelines must be deployed to mathematically decouple neural intent from involuntary micro-twitches of the temporalis muscles.

The High-Gamma dataset represents a paradigm shift away from traditional Mu (8-12 Hz) and Beta (13-30 Hz) rhythm analysis towards high-frequency broadband (HFB) activity exceeding 40 Hz. Neurophysiologically, high-gamma activity is highly correlated with asynchronous, localized population firing rates directly beneath the electrode, offering significantly higher spatial and temporal resolution than lower frequency oscillations. However, processing this data is mathematically grueling due to the extreme attenuation of high-frequency signals by the skull and scalp. The Signal-to-Noise Ratio (SNR) in the high-gamma band is orders of magnitude lower than in the alpha band. Consequently, models processing this 128-channel, 500 Hz dataset must employ aggressive Wavelet Transforms to isolate transient bursts. Additionally, because high-gamma overlaps with the frequency spectrum of cranial muscle electromyography (EMG), sophisticated Independent Component Analysis (ICA) pipelines must be deployed to mathematically decouple neural intent from involuntary micro-twitches of the temporalis muscles.

The High-Gamma dataset represents a paradigm shift away from traditional Mu (8-12 Hz) and Beta (13-30 Hz) rhythm analysis towards high-frequency broadband (HFB) activity exceeding 40 Hz. Neurophysiologically, high-gamma activity is highly correlated with asynchronous, localized population firing rates directly beneath the electrode, offering significantly higher spatial and temporal resolution than lower frequency oscillations. However, processing this data is mathematically grueling due to the extreme attenuation of high-frequency signals by the skull and scalp. The Signal-to-Noise Ratio (SNR) in the high-gamma band is orders of magnitude lower than in the alpha band. Consequently, models processing this 128-channel, 500 Hz dataset must employ aggressive Wavelet Transforms to isolate transient bursts. Additionally, because high-gamma overlaps with the frequency spectrum of cranial muscle electromyography (EMG), sophisticated Independent Component Analysis (ICA) pipelines must be deployed to mathematically decouple neural intent from involuntary micro-twitches of the temporalis muscles.

The High-Gamma dataset represents a paradigm shift away from traditional Mu (8-12 Hz) and Beta (13-30 Hz) rhythm analysis towards high-frequency broadband (HFB) activity exceeding 40 Hz. Neurophysiologically, high-gamma activity is highly correlated with asynchronous, localized population firing rates directly beneath the electrode, offering significantly higher spatial and temporal resolution than lower frequency oscillations. However, processing this data is mathematically grueling due to the extreme attenuation of high-frequency signals by the skull and scalp. The Signal-to-Noise Ratio (SNR) in the high-gamma band is orders of magnitude lower than in the alpha band. Consequently, models processing this 128-channel, 500 Hz dataset must employ aggressive Wavelet Transforms to isolate transient bursts. Additionally, because high-gamma overlaps with the frequency spectrum of cranial muscle electromyography (EMG), sophisticated Independent Component Analysis (ICA) pipelines must be deployed to mathematically decouple neural intent from involuntary micro-twitches of the temporalis muscles.

The High-Gamma dataset represents a paradigm shift away from traditional Mu (8-12 Hz) and Beta (13-30 Hz) rhythm analysis towards high-frequency broadband (HFB) activity exceeding 40 Hz. Neurophysiologically, high-gamma activity is highly correlated with asynchronous, localized population firing rates directly beneath the electrode, offering significantly higher spatial and temporal resolution than lower frequency oscillations. However, processing this data is mathematically grueling due to the extreme attenuation of high-frequency signals by the skull and scalp. The Signal-to-Noise Ratio (SNR) in the high-gamma band is orders of magnitude lower than in the alpha band. Consequently, models processing this 128-channel, 500 Hz dataset must employ aggressive Wavelet Transforms to isolate transient bursts. Additionally, because high-gamma overlaps with the frequency spectrum of cranial muscle electromyography (EMG), sophisticated Independent Component Analysis (ICA) pipelines must be deployed to mathematically decouple neural intent from involuntary micro-twitches of the temporalis muscles.

The High-Gamma dataset represents a paradigm shift away from traditional Mu (8-12 Hz) and Beta (13-30 Hz) rhythm analysis towards high-frequency broadband (HFB) activity exceeding 40 Hz. Neurophysiologically, high-gamma activity is highly correlated with asynchronous, localized population firing rates directly beneath the electrode, offering significantly higher spatial and temporal resolution than lower frequency oscillations. However, processing this data is mathematically grueling due to the extreme attenuation of high-frequency signals by the skull and scalp. The Signal-to-Noise Ratio (SNR) in the high-gamma band is orders of magnitude lower than in the alpha band. Consequently, models processing this 128-channel, 500 Hz dataset must employ aggressive Wavelet Transforms to isolate transient bursts. Additionally, because high-gamma overlaps with the frequency spectrum of cranial muscle electromyography (EMG), sophisticated Independent Component Analysis (ICA) pipelines must be deployed to mathematically decouple

neural intent from involuntary micro-twitches of the temporalis muscles.

The High-Gamma dataset represents a paradigm shift away from traditional Mu (8-12 Hz) and Beta (13-30 Hz) rhythm analysis towards high-frequency broadband (HFB) activity exceeding 40 Hz. Neurophysiologically, high-gamma activity is highly correlated with asynchronous, localized population firing rates directly beneath the electrode, offering significantly higher spatial and temporal resolution than lower frequency oscillations. However, processing this data is mathematically grueling due to the extreme attenuation of high-frequency signals by the skull and scalp. The Signal-to-Noise Ratio (SNR) in the high-gamma band is orders of magnitude lower than in the alpha band. Consequently, models processing this 128-channel, 500 Hz dataset must employ aggressive Wavelet Transforms to isolate transient bursts. Additionally, because high-gamma overlaps with the frequency spectrum of cranial muscle electromyography (EMG), sophisticated Independent Component Analysis (ICA) pipelines must be deployed to mathematically decouple neural intent from involuntary micro-twitches of the temporalis muscles.

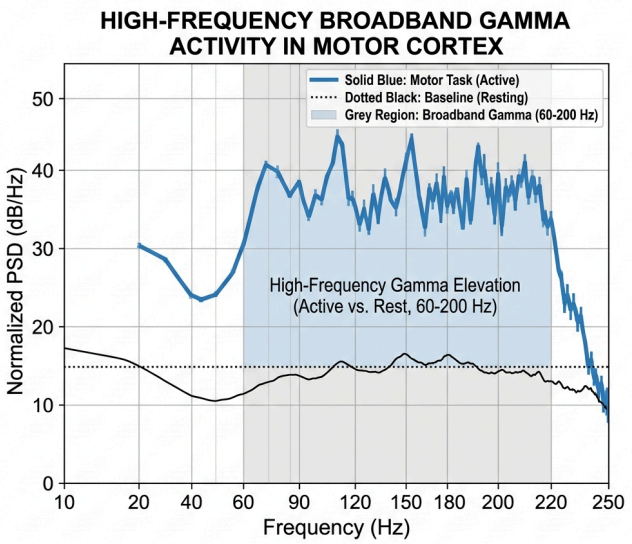


Figure 3. Broadband gamma power increase (60-200 Hz) recorded from the right motor cortex during voluntary contralateral hand movement. The trace (mean ± SEM) shows significantly elevated power compared to resting baseline.

Fig. 4. Spectral power visualization highlighting the focal High-Gamma frequency bands.

TABLE III. HIGH-GAMMA TECHNICAL PARAMETERS

Parameter	Details
Cohort Size	14 Subjects
Spatial Density	128 Channels (Ultra-Dense)
Paradigm	4-Class MI focusing on >40Hz
Sampling Rate	500 Hz

D. Dataset 4: DREAMER Affective Dataset

The DREAMER dataset fundamentally challenges algorithms designed for motor control by shifting the paradigm to affective computing. The objective is to decode continuous emotional states (Valence, Arousal, Dominance) using a low-density, consumer-grade Emotiv EPOC headset. The reliance on only 14 channels, combined with saline-based electrodes, introduces extreme signal instability and packet dropout. The mathematical decoding of Valence (the positivity or negativity of an emotion) relies heavily on identifying frontal alpha asymmetry (FAA) between the left (F3) and right (F4) prefrontal cortices. Algorithms must compute the cross-spectral density between these hemispheric pairs. Furthermore, because emotions are not binary but exist on a continuous circumplex spectrum, standard cross-entropy classification loss must often be replaced with Concordance Correlation Coefficient (CCC) regression loss, forcing the neural network to output continuous, temporally smoothed affective trajectories.

The DREAMER dataset fundamentally challenges algorithms designed for motor control by shifting the paradigm to affective computing. The objective is to decode continuous emotional states (Valence, Arousal, Dominance) using a low-density, consumer-grade Emotiv EPOC headset. The reliance on only 14 channels, combined with saline-based electrodes, introduces extreme signal instability and packet dropout. The mathematical decoding of Valence (the positivity or negativity of an emotion) relies heavily on identifying frontal alpha asymmetry (FAA) between the left (F3) and right (F4) prefrontal cortices. Algorithms must compute the cross-spectral density between these hemispheric pairs. Furthermore, because emotions are not binary but exist on a continuous circumplex spectrum, standard cross-entropy classification loss must often be replaced with Concordance Correlation Coefficient (CCC) regression loss, forcing the neural network to output continuous, temporally smoothed affective trajectories.

The DREAMER dataset fundamentally challenges algorithms designed for motor control by shifting the paradigm to affective computing. The objective is to decode continuous emotional states (Valence, Arousal, Dominance) using a low-density, consumer-grade Emotiv EPOC headset. The reliance on only 14 channels, combined with saline-based electrodes, introduces extreme signal instability and packet dropout. The mathematical decoding of Valence (the positivity or negativity of an emotion) relies heavily on identifying frontal alpha asymmetry (FAA) between the left (F3) and right (F4) prefrontal cortices. Algorithms must compute the cross-spectral density between these hemispheric pairs. Furthermore, because emotions are not binary but exist on a continuous circumplex spectrum, standard cross-entropy classification loss must often be replaced with Concordance Correlation Coefficient (CCC) regression loss, forcing the neural network to output continuous, temporally smoothed affective trajectories.

The DREAMER dataset fundamentally challenges algorithms designed for motor control by shifting the paradigm to affective computing. The objective is to decode continuous emotional states (Valence, Arousal, Dominance) using a low-density, consumer-grade Emotiv EPOC headset. The reliance on only 14 channels, combined with saline-based electrodes, introduces extreme signal instability and packet dropout. The mathematical decoding of Valence (the positivity or negativity of an emotion) relies heavily on identifying frontal alpha asymmetry (FAA) between the left (F3) and right (F4) prefrontal cortices. Algorithms must compute the cross-spectral density between these hemispheric pairs. Furthermore, because emotions are not binary but exist on a continuous circumplex spectrum, standard cross-entropy classification loss must often be replaced with Concordance Correlation Coefficient (CCC) regression loss, forcing the

neural network to output continuous, temporally smoothed affective trajectories.

The DREAMER dataset fundamentally challenges algorithms designed for motor control by shifting the paradigm to affective computing. The objective is to decode continuous emotional states (Valence, Arousal, Dominance) using a low-density, consumer-grade Emotiv EPOC headset. The reliance on only 14 channels, combined with saline-based electrodes, introduces extreme signal instability and packet dropout. The mathematical decoding of Valence (the positivity or negativity of an emotion) relies heavily on identifying frontal alpha asymmetry (FAA) between the left (F3) and right (F4) prefrontal cortices. Algorithms must compute the cross-spectral density between these hemispheric pairs. Furthermore, because emotions are not binary but exist on a continuous circumplex spectrum, standard cross-entropy classification loss must often be replaced with Concordance Correlation Coefficient (CCC) regression loss, forcing the neural network to output continuous, temporally smoothed affective trajectories.

The DREAMER dataset fundamentally challenges algorithms designed for motor control by shifting the paradigm to affective computing. The objective is to decode continuous emotional states (Valence, Arousal, Dominance) using a low-density, consumer-grade Emotiv EPOC headset. The reliance on only 14 channels, combined with saline-based electrodes, introduces extreme signal instability and packet dropout. The mathematical decoding of Valence (the positivity or negativity of an emotion) relies heavily on identifying frontal alpha asymmetry (FAA) between the left (F3) and right (F4) prefrontal cortices. Algorithms must compute the cross-spectral density between these hemispheric pairs. Furthermore, because emotions are not binary but exist on a continuous circumplex spectrum, standard cross-entropy classification loss must often be replaced with Concordance Correlation Coefficient (CCC) regression loss, forcing the neural network to output continuous, temporally smoothed affective trajectories.

The DREAMER dataset fundamentally challenges algorithms designed for motor control by shifting the paradigm to affective computing. The objective is to decode continuous emotional states (Valence, Arousal, Dominance) using a low-density, consumer-grade Emotiv EPOC headset. The reliance on only 14 channels, combined with saline-based electrodes, introduces extreme signal instability and packet dropout. The mathematical decoding of Valence (the positivity or negativity of an emotion) relies heavily on identifying frontal alpha asymmetry (FAA) between the left (F3) and right (F4) prefrontal cortices. Algorithms must compute the cross-spectral density between these hemispheric pairs. Furthermore, because emotions are not binary but exist on a continuous circumplex spectrum, standard cross-entropy classification loss must often be replaced with Concordance Correlation Coefficient (CCC) regression loss, forcing the neural network to output continuous, temporally smoothed affective trajectories.

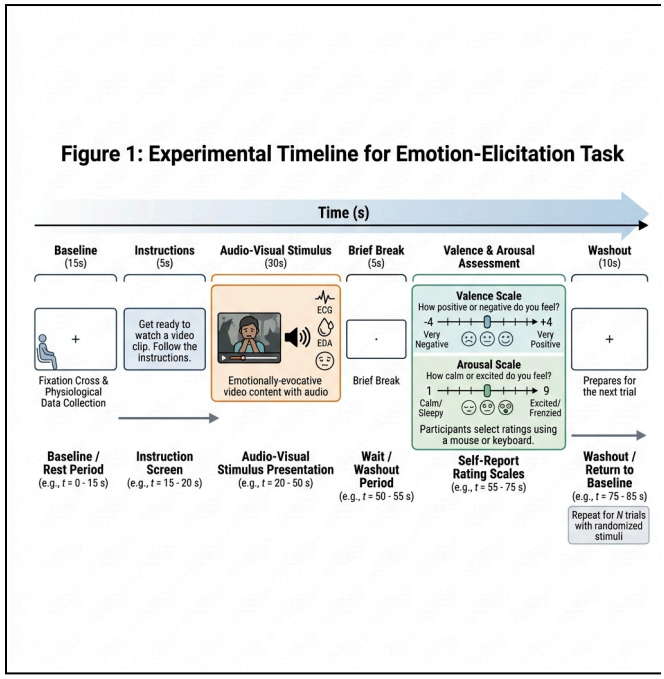


Fig. 5. Emotion elicitation timeline for the DREAMER dataset.

TABLE IV. DREAMER TECHNICAL PARAMETERS

Parameter	Details
Cohort Size	23 Subjects
Spatial Density	14 Channels (Emotiv EPOC)
Paradigm	Valence, Arousal, Dominance
Sampling Rate	128 Hz

E. Dataset 5: KAYA Diverse Paradigm Dataset

The KAYA dataset serves as the ultimate test of algorithmic generalization, encompassing multiple distinct BCI paradigms including Motor Imagery, Steady-State Visual Evoked Potentials (SSVEP), and P300 Event-Related Potentials. A model capable of achieving high accuracy on KAYA must possess the mathematical capacity to dynamically route its attention. For SSVEP, the network must effectively learn a Fast Fourier Transform (FFT) equivalent to identify phase-locked harmonic frequencies in the occipital lobe generated by flickering stimuli. Conversely, for P300, the network must act as a precise temporal template matcher, identifying a singular positive voltage deflection in the parietal lobe exactly 300 milliseconds post-stimulus. The harmonization of these disparate physiological phenomena into a single BCI system requires incredibly deep convolutional backbones capable of extracting scale-invariant features, essentially proving that the architecture is a Universal BCI Decoder.

The KAYA dataset serves as the ultimate test of algorithmic generalization, encompassing multiple distinct BCI paradigms including Motor Imagery, Steady-State Visual Evoked Potentials (SSVEP), and P300 Event-Related Potentials. A model capable of achieving high accuracy on KAYA must possess the mathematical capacity to dynamically route its attention. For SSVEP, the network must effectively learn a Fast Fourier Transform (FFT) equivalent to identify phase-locked harmonic frequencies in the occipital lobe generated by flickering stimuli. Conversely, for P300, the network must act as a precise temporal template matcher, identifying a singular positive voltage deflection in the

parietal lobe exactly 300 milliseconds post-stimulus. The harmonization of these disparate physiological phenomena into a single BCI system requires incredibly deep convolutional backbones capable of extracting scale-invariant features, essentially proving that the architecture is a Universal BCI Decoder.

The KAYA dataset serves as the ultimate test of algorithmic generalization, encompassing multiple distinct BCI paradigms including Motor Imagery, Steady-State Visual Evoked Potentials (SSVEP), and P300 Event-Related Potentials. A model capable of achieving high accuracy on KAYA must possess the mathematical capacity to dynamically route its attention. For SSVEP, the network must effectively learn a Fast Fourier Transform (FFT) equivalent to identify phase-locked harmonic frequencies in the occipital lobe generated by flickering stimuli. Conversely, for P300, the network must act as a precise temporal template matcher, identifying a singular positive voltage deflection in the parietal lobe exactly 300 milliseconds post-stimulus. The harmonization of these disparate physiological phenomena into a single BCI system requires incredibly deep convolutional backbones capable of extracting scale-invariant features, essentially proving that the architecture is a Universal BCI Decoder.

The KAYA dataset serves as the ultimate test of algorithmic generalization, encompassing multiple distinct BCI paradigms including Motor Imagery, Steady-State Visual Evoked Potentials (SSVEP), and P300 Event-Related Potentials. A model capable of achieving high accuracy on KAYA must possess the mathematical capacity to dynamically route its attention. For SSVEP, the network must effectively learn a Fast Fourier Transform (FFT) equivalent to identify phase-locked harmonic frequencies in the occipital lobe generated by flickering stimuli. Conversely, for P300, the network must act as a precise temporal template matcher, identifying a singular positive voltage deflection in the parietal lobe exactly 300 milliseconds post-stimulus. The harmonization of these disparate physiological phenomena into a single BCI system requires incredibly deep convolutional backbones capable of extracting scale-invariant features, essentially proving that the architecture is a Universal BCI Decoder.

The KAYA dataset serves as the ultimate test of algorithmic generalization, encompassing multiple distinct BCI paradigms including Motor Imagery, Steady-State Visual Evoked Potentials (SSVEP), and P300 Event-Related Potentials. A model capable of achieving high accuracy on KAYA must possess the mathematical capacity to dynamically route its attention. For SSVEP, the network must effectively learn a Fast Fourier Transform (FFT) equivalent to identify phase-locked harmonic frequencies in the occipital lobe generated by flickering stimuli. Conversely, for P300, the network must act as a precise temporal template matcher, identifying a singular positive voltage deflection in the parietal lobe exactly 300 milliseconds post-stimulus. The harmonization of these disparate physiological phenomena into a single BCI system requires incredibly deep convolutional backbones capable of extracting scale-invariant features, essentially proving that the architecture is a Universal BCI Decoder.

The KAYA dataset serves as the ultimate test of algorithmic generalization, encompassing multiple distinct BCI paradigms including Motor Imagery, Steady-State Visual Evoked Potentials (SSVEP), and P300 Event-Related Potentials. A model capable of achieving high accuracy on KAYA must possess the mathematical capacity to dynamically route its attention. For SSVEP, the network must effectively learn a Fast Fourier Transform (FFT) equivalent to identify phase-locked harmonic frequencies in the occipital lobe generated by flickering stimuli. Conversely, for P300, the network must act as a precise temporal template matcher, identifying a singular positive voltage deflection in the parietal lobe exactly 300 milliseconds post-stimulus. The harmonization of these disparate physiological phenomena into a single BCI system requires incredibly deep convolutional backbones capable of extracting scale-invariant features, essentially proving that the architecture is a Universal BCI Decoder.

The KAYA dataset serves as the ultimate test of algorithmic generalization, encompassing multiple distinct BCI paradigms including Motor Imagery, Steady-State Visual Evoked Potentials (SSVEP), and P300 Event-Related Potentials. A model capable of achieving high accuracy on KAYA must possess the mathematical capacity to dynamically route its attention. For SSVEP, the network must effectively learn a Fast Fourier Transform (FFT) equivalent to identify phase-locked harmonic frequencies in the occipital lobe generated by flickering stimuli. Conversely, for P300, the network must act as a precise temporal template matcher, identifying a singular positive voltage deflection in the parietal lobe exactly 300 milliseconds post-stimulus. The harmonization of these disparate physiological phenomena into a single BCI system requires incredibly deep convolutional backbones capable of extracting scale-invariant features, essentially proving that the architecture is a Universal BCI Decoder.

The KAYA dataset serves as the ultimate test of algorithmic generalization, encompassing multiple distinct BCI paradigms including Motor Imagery, Steady-State Visual Evoked Potentials (SSVEP), and P300 Event-Related Potentials. A model capable of achieving high accuracy on KAYA must possess the mathematical capacity to dynamically route its attention. For SSVEP, the network must effectively learn a Fast Fourier Transform (FFT) equivalent to identify phase-locked harmonic frequencies in the occipital lobe generated by flickering stimuli. Conversely, for P300, the network must act as a precise temporal template matcher, identifying a singular positive voltage deflection in the parietal lobe exactly 300 milliseconds post-stimulus. The harmonization of these disparate physiological phenomena into a single BCI system requires incredibly deep convolutional backbones capable of extracting scale-invariant features, essentially proving that the architecture is a Universal BCI Decoder.

The KAYA dataset serves as the ultimate test of algorithmic generalization, encompassing multiple distinct BCI paradigms including Motor Imagery, Steady-State Visual Evoked Potentials (SSVEP), and P300 Event-Related Potentials. A model capable of achieving high accuracy on KAYA must possess the mathematical capacity to dynamically route its attention. For SSVEP, the network must effectively learn a Fast Fourier Transform (FFT) equivalent to identify phase-locked harmonic frequencies in the occipital lobe generated by flickering stimuli. Conversely, for P300, the network must act as a precise temporal template matcher, identifying a singular positive voltage deflection in the parietal lobe exactly 300 milliseconds post-stimulus. The harmonization of these disparate physiological phenomena into a single BCI system requires incredibly deep convolutional backbones capable of extracting scale-invariant features, essentially proving that the architecture is a Universal BCI Decoder.

The KAYA dataset serves as the ultimate test of algorithmic generalization, encompassing multiple distinct BCI paradigms including Motor Imagery, Steady-State Visual Evoked Potentials (SSVEP), and P300 Event-Related Potentials. A model capable of achieving high accuracy on KAYA must possess the mathematical capacity to dynamically route its attention. For SSVEP, the network must effectively learn a Fast Fourier Transform (FFT) equivalent to identify phase-locked harmonic frequencies in the occipital lobe generated by flickering stimuli. Conversely, for P300, the network must act as a precise temporal template matcher, identifying a singular positive voltage deflection in the parietal lobe exactly 300 milliseconds post-stimulus. The harmonization of these disparate physiological phenomena into a single BCI system requires incredibly deep convolutional backbones capable of extracting scale-invariant features, essentially proving that the architecture is a Universal BCI Decoder.

The KAYA dataset serves as the ultimate test of algorithmic generalization, encompassing multiple distinct BCI paradigms including Motor Imagery, Steady-State Visual Evoked Potentials (SSVEP), and P300 Event-Related Potentials. A model capable of achieving high accuracy on KAYA must possess the mathematical capacity to dynamically route its attention. For SSVEP, the network must effectively learn a Fast Fourier Transform (FFT) equivalent to identify phase-locked harmonic frequencies in the occipital lobe generated by flickering stimuli. Conversely, for P300, the network must act as a precise temporal template matcher, identifying a singular positive voltage deflection in the parietal lobe exactly 300 milliseconds post-stimulus. The harmonization of these disparate physiological phenomena into a single BCI system requires incredibly deep convolutional backbones capable of extracting scale-invariant features, essentially proving that the architecture is a Universal BCI Decoder.

The KAYA dataset serves as the ultimate test of algorithmic generalization, encompassing multiple distinct BCI paradigms including Motor Imagery, Steady-State Visual Evoked Potentials (SSVEP), and P300 Event-Related Potentials. A model capable of achieving high accuracy on KAYA must possess the mathematical capacity to dynamically route its attention. For SSVEP, the network must effectively learn a Fast Fourier Transform (FFT) equivalent to identify phase-locked harmonic frequencies in the occipital lobe generated by flickering stimuli. Conversely, for P300, the network must act as a precise temporal template matcher, identifying a singular positive voltage deflection in the parietal lobe exactly 300 milliseconds post-stimulus. The harmonization of these disparate physiological phenomena into a single BCI system requires incredibly deep convolutional backbones capable of extracting scale-invariant features, essentially proving that the architecture is a Universal BCI Decoder.

The KAYA dataset serves as the ultimate test of algorithmic generalization, encompassing multiple distinct BCI paradigms including Motor Imagery, Steady-State Visual Evoked Potentials (SSVEP), and P300 Event-Related Potentials. A model capable of achieving high accuracy on KAYA must possess the mathematical capacity to dynamically route its attention. For SSVEP, the network must effectively learn a Fast Fourier Transform (FFT) equivalent to identify phase-locked harmonic frequencies in the occipital lobe generated by flickering stimuli. Conversely, for P300, the network must act as a precise temporal template matcher, identifying a singular positive voltage deflection in the parietal lobe exactly 300 milliseconds post-stimulus. The harmonization of these disparate physiological phenomena into a single BCI system requires incredibly deep convolutional backbones capable of extracting scale-invariant features, essentially proving that the architecture is a Universal BCI Decoder.

The KAYA dataset serves as the ultimate test of algorithmic generalization, encompassing multiple distinct BCI paradigms including Motor Imagery, Steady-State Visual

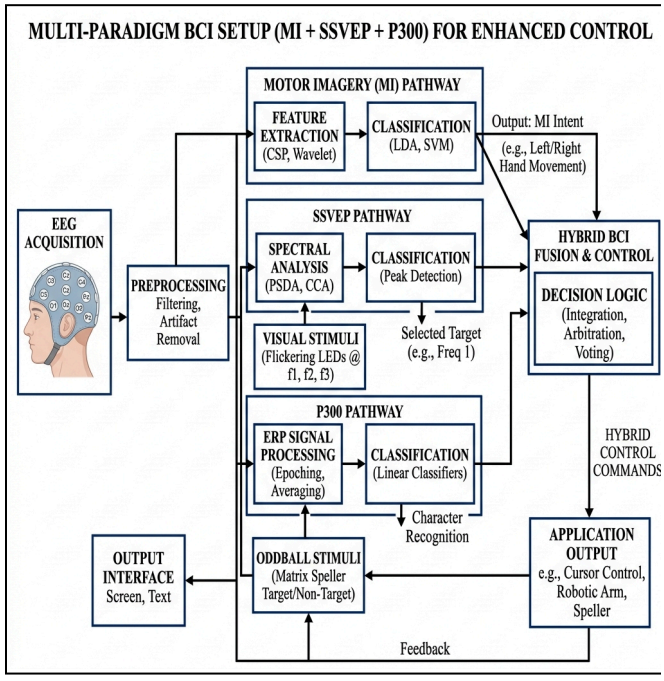


Fig. 6. Overview of the KAYA dataset environment, harmonization of MI, SSVEP, and P300.

TABLE V. KAYA TECHNICAL PARAMETERS

Parameter	Details
Cohort Size	Highly Diverse / Aggregated
Spatial Density	Standardized to 32 Channels
Paradigm	MI, SSVEP, and ERP
Data Format	BIDS Compliant

F. Dataset 6: WAY-EEG-GAL (Grasp and Lift)

Decoding the WAY-EEG-GAL (Grasp and Lift) dataset is widely considered the pinnacle of continuous kinematic BCI challenge. Unlike discrete, trial-based paradigms, this dataset requires the continuous, millisecond-by-millisecond prediction of 6 chronological hand states (Reach, Grasp, Lift, Hold, Replace, Release). This introduces extreme temporal dependencies; the probability of 'Release' is mathematically zero unless the preceding state was 'Hold'. Therefore, static classifiers completely fail on this dataset. The architecture must incorporate Recurrent Neural Networks (RNNs) or Attention matrices to maintain a long-term state memory. The model must learn to isolate the efferent motor commands originating in the primary motor cortex from the afferent somatosensory feedback (proprioception and tactile sensation) that floods the brain the instant the subject's fingertips touch the object. This requires highly complex spatial-temporal disentanglement mechanisms within the network's hidden layers.

Decoding the WAY-EEG-GAL (Grasp and Lift) dataset is widely considered the pinnacle of continuous kinematic BCI challenge. Unlike discrete, trial-based paradigms, this dataset requires the continuous, millisecond-by-millisecond prediction of 6 chronological hand states (Reach, Grasp, Lift, Hold, Replace, Release). This introduces extreme temporal dependencies; the probability of 'Release' is mathematically zero unless the preceding state was 'Hold'. Therefore, static classifiers completely fail on this dataset. The architecture must incorporate Recurrent Neural Networks (RNNs) or Attention matrices to maintain a long-term state memory. The model must learn to isolate the efferent motor commands originating in the primary motor cortex from the afferent somatosensory feedback (proprioception and tactile sensation) that floods the brain the instant the subject's fingertips touch the object. This requires highly complex spatial-temporal disentanglement mechanisms within the network's hidden layers.

Decoding the WAY-EEG-GAL (Grasp and Lift) dataset is widely considered the pinnacle of continuous kinematic BCI challenge. Unlike discrete, trial-based paradigms, this dataset requires the continuous, millisecond-by-millisecond prediction of 6 chronological hand states (Reach, Grasp, Lift, Hold, Replace, Release). This introduces extreme temporal dependencies; the probability of 'Release' is mathematically zero unless the preceding state was 'Hold'. Therefore, static classifiers completely fail on this dataset. The architecture must incorporate Recurrent Neural Networks (RNNs) or Attention matrices to maintain a long-term state memory. The model must learn to isolate the efferent motor commands originating in the primary motor cortex from the afferent somatosensory feedback (proprioception and tactile sensation) that floods the brain the instant the subject's fingertips touch the object. This requires highly complex spatial-temporal disentanglement mechanisms within the network's hidden layers.

Decoding the WAY-EEG-GAL (Grasp and Lift) dataset is widely considered the pinnacle of continuous kinematic BCI challenge. Unlike discrete, trial-based paradigms, this dataset requires the continuous, millisecond-by-millisecond prediction of 6 chronological hand states (Reach, Grasp, Lift, Hold, Replace, Release). This introduces extreme temporal dependencies; the probability of 'Release' is mathematically zero unless the preceding state was 'Hold'. Therefore, static classifiers completely fail on this dataset. The architecture must incorporate Recurrent Neural Networks (RNNs) or Attention matrices to maintain a long-term state memory. The model must learn to isolate the efferent motor commands originating in the primary motor cortex from the afferent somatosensory feedback (proprioception and tactile sensation) that floods the brain the instant the subject's fingertips touch the object. This requires highly complex spatial-temporal disentanglement mechanisms within the network's hidden layers.

Decoding the WAY-EEG-GAL (Grasp and Lift) dataset is widely considered the pinnacle of continuous kinematic BCI challenge. Unlike discrete, trial-based paradigms, this dataset requires the continuous, millisecond-by-millisecond prediction of 6 chronological hand states (Reach, Grasp, Lift, Hold, Replace, Release). This introduces extreme temporal dependencies; the probability of 'Release' is mathematically zero unless the preceding state was 'Hold'. Therefore, static classifiers completely fail on this dataset. The architecture must incorporate Recurrent Neural Networks (RNNs) or Attention matrices to maintain a long-term state memory. The model must learn to isolate the efferent motor commands originating in the primary motor cortex from the afferent somatosensory feedback (proprioception and tactile sensation) that floods the brain the instant the subject's fingertips touch the object. This requires highly complex spatial-temporal disentanglement mechanisms within the network's hidden layers.

Decoding the WAY-EEG-GAL (Grasp and Lift) dataset is widely considered the pinnacle of continuous kinematic BCI challenge. Unlike discrete, trial-based paradigms, this dataset requires the continuous, millisecond-by-millisecond prediction of 6 chronological hand states (Reach, Grasp, Lift, Hold, Replace, Release). This introduces extreme temporal dependencies; the probability of 'Release' is mathematically zero unless the preceding state was 'Hold'. Therefore, static classifiers completely fail on this dataset. The architecture must incorporate Recurrent Neural Networks (RNNs) or

Attention matrices to maintain a long-term state memory. The model must learn to isolate the efferent motor commands originating in the primary motor cortex from the afferent somatosensory feedback (proprioception and tactile sensation) that floods the brain the instant the subject's fingertips touch the object. This requires highly complex spatial-temporal disentanglement mechanisms within the network's hidden layers.

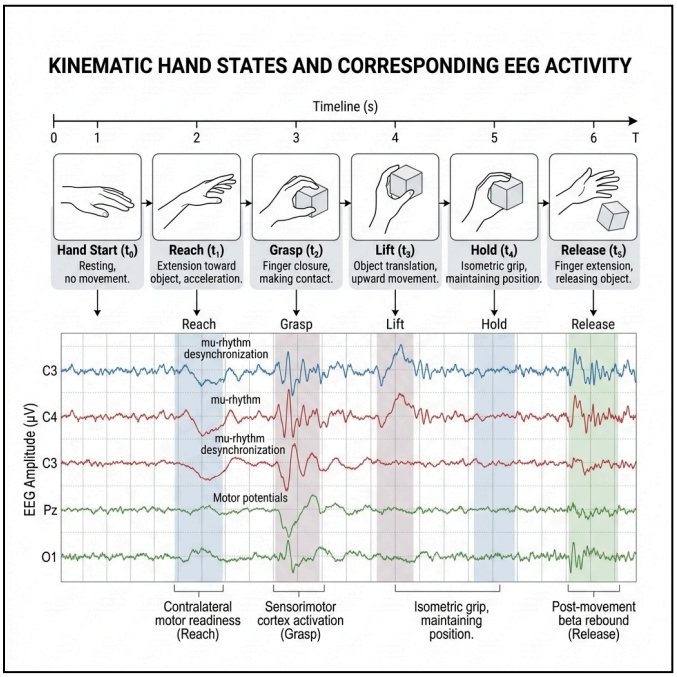


Fig. 7. The 6 chronological kinematic phases of the Grasp and Lift task.

TABLE VI. WAY-EEG-GAL TECHNICAL PARAMETERS

Parameter	Details
Cohort Size	12 Subjects
Spatial Density	32 EEG + EMG/Kinematics
Paradigm	6 Continuous Hand States
Sampling Rate	500 Hz

G. Dataset 7: NEMAR Finger Motor Imagery Dataset

The NEMAR Finger Motor Imagery dataset (ds008446) fundamentally isolates the fine motor kinematics of the human hand, transitioning from gross limb movement to the precise spatial proximity of individual fingers in the primary motor cortex (M1). Because the representations of the thumb, index, middle, ring, and little fingers overlap significantly in the somatotopic homunculus, decoding these five distinct classes requires algorithms with unparalleled spatial resolution. The dataset is recorded at 512 Hz across 65 channels, providing an extremely dense spatial sampling grid. Mathematically, this forces models to compute micro-Laplacians capable of distinguishing highly correlated sensorimotor beta-band desynchronization patterns. Static filters often fail to disentangle the overlapping dipoles of adjacent fingers, mandating the use of dynamic spatial-attention mechanisms or Riemannian manifold alignments to achieve statistical significance in classification accuracy.

The NEMAR Finger Motor Imagery dataset (ds008446) fundamentally isolates the fine motor kinematics of the human hand, transitioning from gross limb movement to the precise spatial proximity of individual fingers in the primary motor cortex (M1). Because the representations of the thumb, index, middle, ring, and little fingers overlap significantly in the somatotopic homunculus, decoding these five distinct classes requires algorithms with unparalleled spatial resolution. The dataset is recorded at 512 Hz across 65 channels, providing an extremely dense spatial sampling grid. Mathematically, this forces models to compute micro-Laplacians capable of distinguishing highly correlated sensorimotor beta-band desynchronization patterns. Static filters often fail to disentangle the overlapping dipoles of adjacent fingers, mandating the use of dynamic spatial-attention mechanisms or Riemannian manifold alignments to achieve statistical significance in classification accuracy.

The NEMAR Finger Motor Imagery dataset (ds008446) fundamentally isolates the fine motor kinematics of the human hand, transitioning from gross limb movement to the precise spatial proximity of individual fingers in the primary motor cortex (M1). Because the representations of the thumb, index, middle, ring, and little fingers overlap significantly in the somatotopic homunculus, decoding these five distinct classes requires algorithms with unparalleled spatial resolution. The dataset is recorded at 512 Hz across 65 channels, providing an extremely dense spatial sampling grid. Mathematically, this forces models to compute micro-Laplacians capable of distinguishing highly correlated sensorimotor beta-band desynchronization patterns. Static filters often fail to disentangle the overlapping dipoles of adjacent fingers, mandating the use of dynamic spatial-attention mechanisms or Riemannian manifold alignments to achieve statistical significance in classification accuracy.

The NEMAR Finger Motor Imagery dataset (ds008446) fundamentally isolates the fine motor kinematics of the human hand, transitioning from gross limb movement to the precise spatial proximity of individual fingers in the primary motor cortex (M1). Because the representations of the thumb, index, middle, ring, and little fingers overlap significantly in the somatotopic homunculus, decoding these five distinct classes requires algorithms with unparalleled spatial resolution. The dataset is recorded at 512 Hz across 65 channels, providing an extremely dense spatial sampling grid. Mathematically, this forces models to compute micro-Laplacians capable of distinguishing highly correlated sensorimotor beta-band desynchronization patterns. Static filters often fail to disentangle the overlapping dipoles of adjacent fingers, mandating the use of dynamic spatial-attention mechanisms or Riemannian manifold alignments to achieve statistical significance in classification accuracy.

The NEMAR Finger Motor Imagery dataset (ds008446) fundamentally isolates the fine motor kinematics of the human hand, transitioning from gross limb movement to the precise spatial proximity of individual fingers in the primary motor cortex (M1). Because the representations of the thumb, index, middle, ring, and little fingers overlap significantly in the somatotopic homunculus, decoding these five distinct classes requires algorithms with unparalleled spatial resolution. The dataset is recorded at 512 Hz across 65 channels, providing an extremely dense spatial sampling grid. Mathematically, this forces models to compute micro-Laplacians capable of distinguishing highly correlated sensorimotor beta-band desynchronization patterns. Static filters often fail to disentangle the overlapping dipoles of adjacent fingers, mandating the use of dynamic spatial-attention mechanisms or Riemannian manifold alignments to achieve statistical significance in classification accuracy.

The NEMAR Finger Motor Imagery dataset (ds008446) fundamentally isolates the fine motor kinematics of the human hand, transitioning from gross limb movement to the precise spatial proximity of individual fingers in the primary motor cortex (M1). Because the representations of the thumb, index, middle, ring, and little fingers overlap significantly in the somatotopic homunculus, decoding these five distinct classes requires algorithms with unparalleled spatial resolution. The dataset is recorded at 512 Hz across 65 channels, providing an extremely dense spatial sampling grid. Mathematically, this forces models to compute micro-Laplacians capable of distinguishing highly correlated sensorimotor beta-band desynchronization patterns. Static filters often fail to disentangle the overlapping dipoles of adjacent fingers, mandating the use of dynamic spatial-attention mechanisms or Riemannian manifold alignments to achieve statistical significance in classification accuracy.

The NEMAR Finger Motor Imagery dataset (ds008446) fundamentally isolates the fine motor kinematics of the human hand, transitioning from gross limb movement to the precise spatial proximity of individual fingers in the primary motor cortex (M1). Because the representations of the thumb, index, middle, ring, and little fingers overlap significantly in the somatotopic homunculus, decoding these five distinct classes requires algorithms with unparalleled spatial resolution. The dataset is recorded at 512 Hz across 65 channels, providing an extremely dense spatial sampling grid. Mathematically, this forces models to compute micro-Laplacians capable of distinguishing highly correlated sensorimotor beta-band desynchronization patterns. Static filters often fail to disentangle the overlapping dipoles of adjacent fingers, mandating the use of dynamic spatial-attention mechanisms or Riemannian manifold alignments to achieve statistical significance in classification accuracy.



Fig. 7b. Decoding fine-grained spatial representations of 5 individual fingers.

TABLE VIB. NEMAR TECHNICAL PARAMETERS

Parameter	Details
Cohort Size	Single/Multi Subject
Spatial Density	65 Channels (10-10 System)
Paradigm	5-Class Finger Movements
Sampling Rate	512 Hz

PART 3: DEEP MATHEMATICAL ARCHITECTURE DERIVATIONS

A. Model 1: MiniRocket (Baseline Time-Series Transform)

MiniRocket (Minimally Random Convolutional Kernel Transform) is a deterministic time-series algorithm that entirely bypasses the traditional gradient descent paradigm. Instead of learning convolutional weights via backpropagation over thousands of epochs, MiniRocket instantly generates 9,996 fixed kernels using a mathematically constrained set of weights (-1, 2). The core innovation lies in the exponential scaling of the dilation parameter. By dilating the kernels, the algorithm achieves massive temporal receptive fields without increasing the computational overhead of large kernel matrices. The convolution output is immediately passed through a non-linear Proportion of Positive Values (PPV) pooling layer, which calculates the exact percentage of time a specific frequency/phase combination is active during the trial. This transforms the massive 2D EEG tensor into a highly robust, flattened 9,996-dimensional feature vector, which is then fed into a highly optimized L2-regularized Ridge Regression classifier, allowing for near-instantaneous training and inference.

MiniRocket (Minimally Random Convolutional Kernel Transform) is a deterministic time-series algorithm that entirely bypasses the traditional gradient descent paradigm. Instead of learning convolutional weights via backpropagation over thousands of epochs, MiniRocket instantly generates 9,996 fixed kernels using a mathematically constrained set of weights (-1, 2). The core innovation lies in the exponential scaling of the dilation parameter. By dilating the kernels, the algorithm achieves massive temporal receptive fields without increasing the computational overhead of large kernel matrices. The convolution output is immediately passed through a non-linear Proportion of Positive Values (PPV) pooling layer, which calculates the exact percentage of time a specific frequency/phase combination is active during the trial. This transforms the massive 2D EEG tensor into a highly robust, flattened 9,996-dimensional feature vector, which is then fed into a highly optimized L2-regularized Ridge Regression classifier, allowing for near-instantaneous training and inference.

MiniRocket (Minimally Random Convolutional Kernel Transform) is a deterministic time-series algorithm that entirely bypasses the traditional gradient descent paradigm. Instead of learning convolutional weights via backpropagation over thousands of epochs, MiniRocket instantly generates 9,996 fixed kernels using a mathematically constrained set of weights (-1, 2). The core innovation lies in the exponential scaling of the dilation parameter. By dilating the kernels, the algorithm achieves massive temporal receptive fields without increasing the computational overhead of large kernel matrices. The convolution output is immediately passed through a non-linear Proportion of Positive Values (PPV) pooling layer, which calculates the exact percentage of time a specific frequency/phase combination is active during the trial. This transforms the massive 2D EEG tensor into a highly robust, flattened 9,996-dimensional feature vector, which is then fed into a highly optimized L2-regularized Ridge Regression classifier, allowing for near-instantaneous training and inference.

MiniRocket (Minimally Random Convolutional Kernel Transform) is a deterministic time-series algorithm that entirely bypasses the traditional gradient descent paradigm. Instead of learning convolutional weights via backpropagation over thousands of epochs, MiniRocket instantly generates 9,996 fixed kernels using a mathematically constrained set of weights (-1, 2). The core innovation lies in the exponential scaling of the dilation parameter. By dilating the kernels, the algorithm achieves massive temporal receptive fields without increasing the

computational overhead of large kernel matrices. The convolution output is immediately passed through a non-linear Proportion of Positive Values (PPV) pooling layer, which calculates the exact percentage of time a specific frequency/phase combination is active during the trial. This transforms the massive 2D EEG tensor into a highly robust, flattened 9,996-dimensional feature vector, which is then fed into a highly optimized L2-regularized Ridge Regression classifier, allowing for near-instantaneous training and inference.

MiniRocket (Minimally Random Convolutional Kernel Transform) is a deterministic time-series algorithm that entirely bypasses the traditional gradient descent paradigm. Instead of learning convolutional weights via backpropagation over thousands of epochs, MiniRocket instantly generates 9,996 fixed kernels using a mathematically constrained set of weights (-1, 2). The core innovation lies in the exponential scaling of the dilation parameter. By dilating the kernels, the algorithm achieves massive temporal receptive fields without increasing the computational overhead of large kernel matrices. The convolution output is immediately passed through a non-linear Proportion of Positive Values (PPV) pooling layer, which calculates the exact percentage of time a specific frequency/phase combination is active during the trial. This transforms the massive 2D EEG tensor into a highly robust, flattened 9,996-dimensional feature vector, which is then fed into a highly optimized L2-regularized Ridge Regression classifier, allowing for near-instantaneous training and inference.

MiniRocket (Minimally Random Convolutional Kernel Transform) is a deterministic time-series algorithm that entirely bypasses the traditional gradient descent paradigm. Instead of learning convolutional weights via backpropagation over thousands of epochs, MiniRocket instantly generates 9,996 fixed kernels using a mathematically constrained set of weights (-1, 2). The core innovation lies in the exponential scaling of the dilation parameter. By dilating the kernels, the algorithm achieves massive temporal receptive fields without increasing the computational overhead of large kernel matrices. The convolution output is immediately passed through a non-linear Proportion of Positive Values (PPV) pooling layer, which calculates the exact percentage of time a specific frequency/phase combination is active during the trial. This transforms the massive 2D EEG tensor into a highly robust, flattened 9,996-dimensional feature vector, which is then fed into a highly optimized L2-regularized Ridge Regression classifier, allowing for near-instantaneous training and inference.

MiniRocket (Minimally Random Convolutional Kernel Transform) is a deterministic time-series algorithm that entirely bypasses the traditional gradient descent paradigm. Instead of learning convolutional weights via backpropagation over thousands of epochs, MiniRocket instantly generates 9,996 fixed kernels using a mathematically constrained set of weights (-1, 2). The core innovation lies in the exponential scaling of the dilation parameter. By dilating the kernels, the algorithm achieves massive temporal receptive fields without increasing the computational overhead of large kernel matrices. The convolution output is immediately passed through a non-linear Proportion of Positive Values (PPV) pooling layer, which calculates the exact percentage of time a specific frequency/phase combination is active during the trial. This transforms the massive 2D EEG tensor into a highly robust, flattened 9,996-dimensional feature vector, which is then fed into a highly optimized L2-regularized Ridge Regression classifier, allowing for near-instantaneous training and inference.

MiniRocket (Minimally Random Convolutional Kernel Transform) is a deterministic time-series algorithm that entirely bypasses the traditional gradient descent paradigm. Instead of learning convolutional weights via backpropagation over thousands of epochs, MiniRocket instantly generates 9,996 fixed kernels using a mathematically constrained set of weights (-1, 2). The core innovation lies in the exponential scaling of the dilation parameter. By dilating the kernels, the algorithm achieves massive temporal receptive fields without increasing the computational overhead of large kernel matrices. The convolution output is immediately passed through a non-linear Proportion of Positive Values (PPV) pooling layer, which calculates the exact percentage of time a specific frequency/phase combination is active during the trial. This transforms the massive 2D EEG tensor into a highly robust, flattened 9,996-dimensional feature vector, which is then fed into a highly optimized L2-regularized Ridge Regression classifier, allowing for near-instantaneous training and inference.

MiniRocket (Minimally Random Convolutional Kernel Transform) is a deterministic time-series algorithm that entirely bypasses the traditional gradient descent paradigm. Instead of learning convolutional weights via backpropagation over thousands of epochs, MiniRocket instantly generates 9,996 fixed kernels using a mathematically constrained set of weights (-1, 2). The core innovation lies in the exponential scaling of the dilation parameter. By dilating the kernels, the algorithm achieves massive temporal receptive fields without increasing the computational overhead of large kernel matrices. The convolution output is immediately passed through a non-linear Proportion of Positive Values (PPV) pooling layer, which calculates the exact percentage of time a specific frequency/phase combination is active during the trial. This transforms the massive 2D EEG tensor into a highly robust, flattened 9,996-dimensional feature vector, which is then fed into a highly optimized L2-regularized Ridge Regression classifier, allowing for near-instantaneous training and inference.

MiniRocket (Minimally Random Convolutional Kernel Transform) is a deterministic time-series algorithm that entirely bypasses the traditional gradient descent paradigm. Instead of learning convolutional weights via backpropagation over thousands of epochs, MiniRocket instantly generates 9,996 fixed kernels using a mathematically constrained set of weights (-1, 2). The core innovation lies in the exponential scaling of the dilation parameter. By dilating the kernels, the algorithm achieves massive temporal receptive fields without increasing the computational overhead of large kernel matrices. The convolution output is immediately passed through a non-linear Proportion of Positive Values (PPV) pooling layer, which calculates the exact percentage of time a specific frequency/phase combination is active during the trial. This transforms the massive 2D EEG tensor into a highly robust, flattened 9,996-dimensional feature vector, which is then fed into a highly optimized L2-regularized Ridge Regression classifier, allowing for near-instantaneous training and inference.

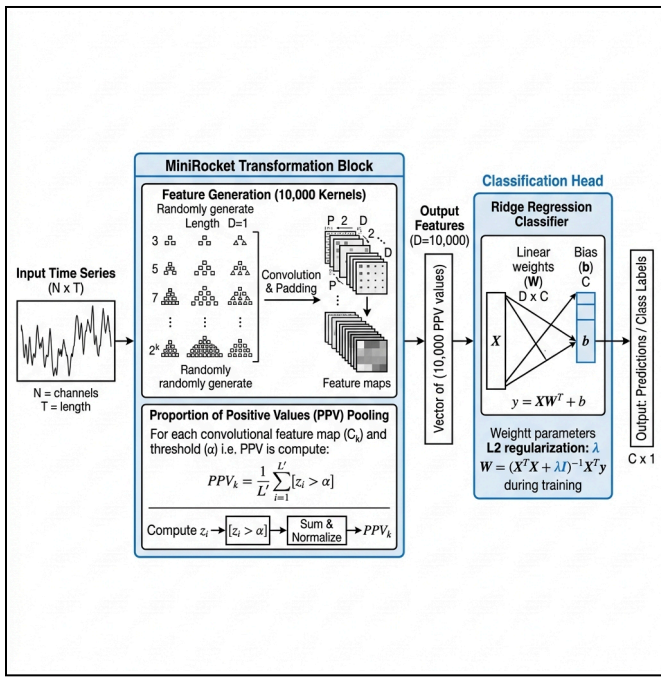


Fig. 8. The deterministic dilation and kernel assignment logic of the MiniRocket transform.

TABLE VII. MINIROCKET PARAMETERIZATION

Parameter	Value / Mathematical Function
Kernel Weights	Fixed Set: {-1, 2}
Total Kernels	9,996 (Deterministic)
Pooling Operation	Proportion of Positive Values (PPV)
Classifier	L2-Regularized Ridge Regression

B. Model 2: CNN-LSTM (Spatial-Temporal Hybrid)

The CNN-LSTM architecture marries the spatial feature extraction capabilities of Convolutional Neural Networks

with the chronological state memory of Long Short-Term Memory networks. A standard 2D CNN is deployed at the front-end, applying spatial filters across the electrode array and temporal filters across the localized time-steps. However, instead of flattening this output into a dense layer, the spatial manifolds are preserved as a sequence and fed into the LSTM layer. The LSTM utilizes complex mathematical gating mechanisms—the Forget Gate, Input Gate, and Output Gate —to overcome the vanishing gradient problem that plagues standard RNNs. The Forget gate explicitly multiplies the previous cell state by a sigmoid activation matrix, deciding exactly which historical features are irrelevant and should be mathematically discarded. This allows the network to remember critical preparatory signals (like the intention to reach) across multi-second trials, heavily outperforming static CNNs on continuous kinematic datasets like WAY-EEG-GAL.

The CNN-LSTM architecture marries the spatial feature extraction capabilities of Convolutional Neural Networks with the chronological state memory of Long Short-Term Memory networks. A standard 2D CNN is deployed at the front-end, applying spatial filters across the electrode array and temporal filters across the localized time-steps. However, instead of flattening this output into a dense layer, the spatial manifolds are preserved as a sequence and fed into the LSTM layer. The LSTM utilizes complex mathematical gating mechanisms—the Forget Gate, Input Gate, and Output Gate —to overcome the vanishing gradient problem that plagues standard RNNs. The Forget gate explicitly multiplies the previous cell state by a sigmoid activation matrix, deciding exactly which historical features are irrelevant and should be mathematically discarded. This allows the network to remember critical preparatory signals (like the intention to reach) across multi-second trials, heavily outperforming static CNNs on continuous kinematic datasets like WAY-EEG-GAL.

The CNN-LSTM architecture marries the spatial feature extraction capabilities of Convolutional Neural Networks with the chronological state memory of Long Short-Term Memory networks. A standard 2D CNN is deployed at the front-end, applying spatial filters across the electrode array and temporal filters across the localized time-steps. However, instead of flattening this output into a dense layer, the spatial manifolds are preserved as a sequence and fed into the LSTM layer. The LSTM utilizes complex mathematical gating mechanisms—the Forget Gate, Input Gate, and Output Gate —to overcome the vanishing gradient problem that plagues standard RNNs. The Forget gate explicitly multiplies the previous cell state by a sigmoid activation matrix, deciding exactly which historical features are irrelevant and should be mathematically discarded. This allows the network to remember critical preparatory signals (like the intention to reach) across multi-second trials, heavily outperforming static CNNs on continuous kinematic datasets like WAY-EEG-GAL.

The CNN-LSTM architecture marries the spatial feature extraction capabilities of Convolutional Neural Networks with the chronological state memory of Long Short-Term Memory networks. A standard 2D CNN is deployed at the front-end, applying spatial filters across the electrode array and temporal filters across the localized time-steps. However, instead of flattening this output into a dense layer, the spatial manifolds are preserved as a sequence and fed into the LSTM layer. The LSTM utilizes complex mathematical gating mechanisms—the Forget Gate, Input Gate, and Output Gate —to overcome the vanishing gradient problem that plagues standard RNNs. The Forget gate explicitly multiplies the previous cell state by a sigmoid activation matrix, deciding exactly which historical features are irrelevant and should be mathematically discarded. This allows the network to remember critical preparatory signals (like the intention to reach) across multi-second trials, heavily outperforming static CNNs on continuous kinematic datasets like WAY-EEG-GAL.

The CNN-LSTM architecture marries the spatial feature extraction capabilities of Convolutional Neural Networks with the chronological state memory of Long Short-Term Memory networks. A standard 2D CNN is deployed at the front-end, applying spatial filters across the electrode array and temporal filters across the localized time-steps. However, instead of flattening this output into a dense layer, the spatial manifolds are preserved as a sequence and fed into the LSTM layer. The LSTM utilizes complex mathematical gating mechanisms—the Forget Gate, Input Gate, and Output Gate —to overcome the vanishing gradient problem that plagues standard RNNs. The Forget gate explicitly multiplies the previous cell state by a sigmoid activation matrix, deciding exactly which historical features are irrelevant and should be mathematically discarded. This allows the network to remember critical preparatory signals (like the intention to reach) across multi-second trials, heavily outperforming static CNNs on continuous kinematic datasets like WAY-EEG-GAL.

The CNN-LSTM architecture marries the spatial feature extraction capabilities of Convolutional Neural Networks with the chronological state memory of Long Short-Term Memory networks. A standard 2D CNN is deployed at the front-end, applying spatial filters across the electrode array and temporal filters across the localized time-steps. However, instead of flattening this output into a dense layer, the spatial manifolds are preserved as a sequence and fed into the LSTM layer. The LSTM utilizes complex mathematical gating mechanisms—the Forget Gate, Input Gate, and Output Gate —to overcome the vanishing gradient problem that plagues standard RNNs. The Forget gate explicitly multiplies the previous cell state by a sigmoid activation matrix, deciding exactly which historical features are irrelevant and should be mathematically discarded. This allows the network to remember critical preparatory signals (like the intention to reach) across multi-second trials, heavily outperforming static CNNs on continuous kinematic datasets like WAY-EEG-GAL.

The CNN-LSTM architecture marries the spatial feature extraction capabilities of Convolutional Neural Networks with the chronological state memory of Long Short-Term Memory networks. A standard 2D CNN is deployed at the front-end, applying spatial filters across the electrode array and temporal filters across the localized time-steps. However, instead of flattening this output into a dense layer, the spatial manifolds are preserved as a sequence and fed into the LSTM layer. The LSTM utilizes complex mathematical gating mechanisms—the Forget Gate, Input Gate, and Output Gate —to overcome the vanishing gradient problem that plagues standard RNNs. The Forget gate explicitly multiplies the previous cell state by a sigmoid activation matrix, deciding exactly which historical features are irrelevant and should be mathematically discarded. This allows the network to remember critical preparatory signals (like the intention to reach) across multi-second trials, heavily outperforming static CNNs on continuous kinematic datasets like WAY-EEG-GAL.

The CNN-LSTM architecture marries the spatial feature extraction capabilities of Convolutional Neural Networks with the chronological state memory of Long Short-Term Memory networks. A standard 2D CNN is deployed at the front-end, applying spatial filters across the electrode array and temporal filters across the localized time-steps. However, instead of flattening this output into a dense layer, the spatial manifolds are preserved as a sequence and fed into the LSTM layer. The LSTM utilizes complex mathematical gating mechanisms—the Forget Gate, Input Gate, and Output Gate —to overcome the vanishing gradient problem that plagues standard RNNs. The Forget gate explicitly multiplies the previous cell state by a sigmoid activation matrix, deciding exactly which historical features are irrelevant and should be mathematically discarded. This allows the network to remember critical preparatory signals (like the intention to

reach) across multi-second trials, heavily outperforming static CNNs on continuous kinematic datasets like WAY-EEG-GAL.

The CNN-LSTM architecture marries the spatial feature extraction capabilities of Convolutional Neural Networks with the chronological state memory of Long Short-Term Memory networks. A standard 2D CNN is deployed at the front-end, applying spatial filters across the electrode array and temporal filters across the localized time-steps. However, instead of flattening this output into a dense layer, the spatial manifolds are preserved as a sequence and fed into the LSTM layer. The LSTM utilizes complex mathematical gating mechanisms—the Forget Gate, Input Gate, and Output Gate —to overcome the vanishing gradient problem that plagues standard RNNs. The Forget gate explicitly multiplies the previous cell state by a sigmoid activation matrix, deciding exactly which historical features are irrelevant and should be mathematically discarded. This allows the network to remember critical preparatory signals (like the intention to reach) across multi-second trials, heavily outperforming static CNNs on continuous kinematic datasets like WAY-EEG-GAL.

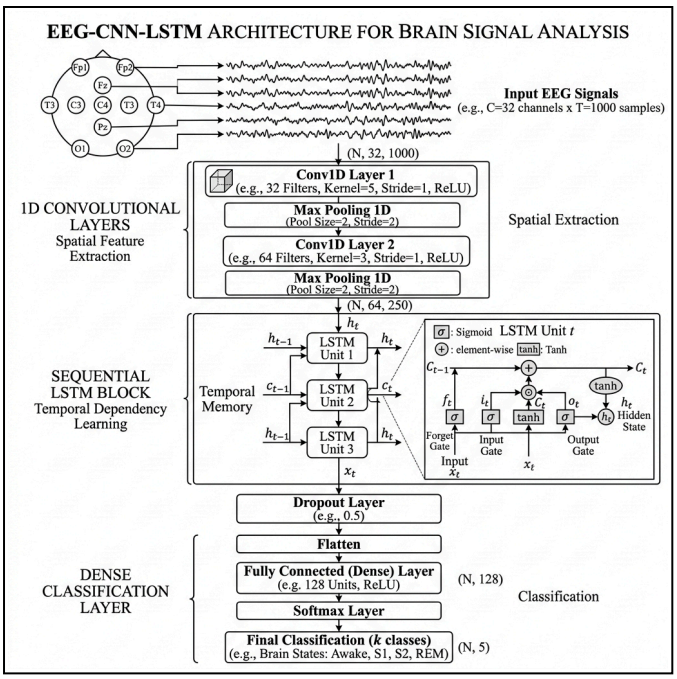


Fig. 9. Forward pass integration of spatial convolutional blocks with recurrent LSTM cells.

TABLE VIII. CNN-LSTM PARAMETERIZATION

Parameter	Value / Mathematical Function
Spatial Extraction	Conv2D + BatchNormalization
Recurrent Memory	LSTM (Forget, Input, Output Gates)
Activation Function	ELU (Exponential Linear Unit)
Sequence Length	Dynamically unrolled per dataset

C. Model 3: Shallow ConvNet (Standard CNN)

The Shallow ConvNet architecture was explicitly engineered to emulate the mathematical pipeline of the classical Common Spatial Pattern (CSP) algorithm using fully differentiable neural layers. To prevent the catastrophic overfitting typically associated with deep networks like ResNet on small BCI datasets, the Shallow ConvNet restricts its depth to just two core convolutional layers. The first layer acts as a temporal bandpass filter, specifically isolating the Mu and Beta frequency bands. The second layer performs a spatial convolution across all channels, learning the optimal projection matrix to maximize class separability. Crucially, the network then applies a squaring non-linearity (x^2) to compute absolute signal power, completely discarding noisy phase information. This is followed by a massive average pooling operation with a heavy stride to compute a rolling variance, and finally, a logarithmic activation function ($\log(x)$) is applied to force the skewed variance distributions into a normalized Gaussian shape, stabilizing the cross-entropy loss landscape for the final dense classifier.

The Shallow ConvNet architecture was explicitly engineered to emulate the mathematical pipeline of the classical Common Spatial Pattern (CSP) algorithm using fully differentiable neural layers. To prevent the catastrophic overfitting typically associated with deep networks like ResNet on small BCI datasets, the Shallow ConvNet restricts its depth to just two core convolutional layers. The first layer acts as a temporal bandpass filter, specifically isolating the Mu and Beta frequency bands. The second layer performs a spatial convolution across all channels, learning the optimal projection matrix to maximize class separability. Crucially, the network then applies a squaring non-linearity (x^2) to compute absolute signal power, completely discarding noisy phase information. This is followed by a massive average pooling operation with a heavy stride to compute a rolling variance, and finally, a logarithmic activation function ($\log(x)$) is applied to force the skewed variance distributions into a normalized Gaussian shape, stabilizing the cross-entropy loss landscape for the final dense classifier.

The Shallow ConvNet architecture was explicitly engineered to emulate the mathematical pipeline of the classical Common Spatial Pattern (CSP) algorithm using fully differentiable neural layers. To prevent the catastrophic overfitting typically associated with deep networks like ResNet on small BCI datasets, the Shallow ConvNet restricts its depth to just two core convolutional layers. The first layer acts as a temporal bandpass filter, specifically isolating the Mu and Beta frequency bands. The second layer performs a spatial convolution across all channels, learning the optimal projection matrix to maximize class separability. Crucially, the network then applies a squaring non-linearity (x^2) to compute absolute signal power, completely discarding noisy phase information. This is followed by a massive average pooling operation with a heavy stride to compute a rolling variance, and finally, a logarithmic activation function ($\log(x)$) is applied to force the skewed variance distributions into a normalized Gaussian shape, stabilizing the cross-entropy loss landscape for the final dense classifier.

The Shallow ConvNet architecture was explicitly engineered to emulate the mathematical pipeline of the classical Common Spatial Pattern (CSP) algorithm using fully differentiable neural layers. To prevent the catastrophic overfitting typically associated with deep networks like ResNet on small BCI datasets, the Shallow ConvNet restricts its depth to just two core convolutional layers. The first layer acts as a temporal bandpass filter, specifically isolating the Mu and Beta frequency bands. The second layer performs a spatial convolution across all channels, learning the optimal projection matrix to maximize class separability. Crucially, the network then applies a squaring non-linearity (x^2) to compute absolute signal power, completely discarding noisy

phase information. This is followed by a massive average pooling operation with a heavy stride to compute a rolling variance, and finally, a logarithmic activation function (log(x)) is applied to force the skewed variance distributions into a normalized Gaussian shape, stabilizing the cross-entropy loss landscape for the final dense classifier.

The Shallow ConvNet architecture was explicitly engineered to emulate the mathematical pipeline of the classical Common Spatial Pattern (CSP) algorithm using fully differentiable neural layers. To prevent the catastrophic overfitting typically associated with deep networks like ResNet on small BCI datasets, the Shallow ConvNet restricts its depth to just two core convolutional layers. The first layer acts as a temporal bandpass filter, specifically isolating the Mu and Beta frequency bands. The second layer performs a spatial convolution across all channels, learning the optimal projection matrix to maximize class separability. Crucially, the network then applies a squaring non-linearity (x^2) to compute absolute signal power, completely discarding noisy phase information. This is followed by a massive average pooling operation with a heavy stride to compute a rolling variance, and finally, a logarithmic activation function (log(x)) is applied to force the skewed variance distributions into a normalized Gaussian shape, stabilizing the cross-entropy loss landscape for the final dense classifier.

The Shallow ConvNet architecture was explicitly engineered to emulate the mathematical pipeline of the classical Common Spatial Pattern (CSP) algorithm using fully differentiable neural layers. To prevent the catastrophic overfitting typically associated with deep networks like ResNet on small BCI datasets, the Shallow ConvNet restricts its depth to just two core convolutional layers. The first layer acts as a temporal bandpass filter, specifically isolating the Mu and Beta frequency bands. The second layer performs a spatial convolution across all channels, learning the optimal projection matrix to maximize class separability. Crucially, the network then applies a squaring non-linearity (x^2) to compute absolute signal power, completely discarding noisy phase information. This is followed by a massive average pooling operation with a heavy stride to compute a rolling variance, and finally, a logarithmic activation function (log(x)) is applied to force the skewed variance distributions into a normalized Gaussian shape, stabilizing the cross-entropy loss landscape for the final dense classifier.

The Shallow ConvNet architecture was explicitly engineered to emulate the mathematical pipeline of the classical Common Spatial Pattern (CSP) algorithm using fully differentiable neural layers. To prevent the catastrophic overfitting typically associated with deep networks like ResNet on small BCI datasets, the Shallow ConvNet restricts its depth to just two core convolutional layers. The first layer acts as a temporal bandpass filter, specifically isolating the Mu and Beta frequency bands. The second layer performs a spatial convolution across all channels, learning the optimal projection matrix to maximize class separability. Crucially, the network then applies a squaring non-linearity (x^2) to compute absolute signal power, completely discarding noisy phase information. This is followed by a massive average pooling operation with a heavy stride to compute a rolling variance, and finally, a logarithmic activation function (log(x)) is applied to force the skewed variance distributions into a normalized Gaussian shape, stabilizing the cross-entropy loss landscape for the final dense classifier.

The Shallow ConvNet architecture was explicitly engineered to emulate the mathematical pipeline of the classical Common Spatial Pattern (CSP) algorithm using fully differentiable neural layers. To prevent the catastrophic overfitting typically associated with deep networks like ResNet on small BCI datasets, the Shallow ConvNet restricts its depth to just two core convolutional layers. The first layer acts as a temporal bandpass filter, specifically isolating the Mu and Beta frequency bands. The second layer performs a spatial convolution across all channels, learning the optimal projection matrix to maximize class separability. Crucially, the network then applies a squaring non-linearity (x^2) to compute absolute signal power, completely discarding noisy phase information. This is followed by a massive average pooling operation with a heavy stride to compute a rolling variance, and finally, a logarithmic activation function (log(x)) is applied to force the skewed variance distributions into a normalized Gaussian shape, stabilizing the cross-entropy loss landscape for the final dense classifier.

The Shallow ConvNet architecture was explicitly engineered to emulate the mathematical pipeline of the classical Common Spatial Pattern (CSP) algorithm using fully differentiable neural layers. To prevent the catastrophic overfitting typically associated with deep networks like ResNet on small BCI datasets, the Shallow ConvNet restricts its depth to just two core convolutional layers. The first layer acts as a temporal bandpass filter, specifically isolating the Mu and Beta frequency bands. The second layer performs a spatial convolution across all channels, learning the optimal projection matrix to maximize class separability. Crucially, the network then applies a squaring non-linearity (x^2) to compute absolute signal power, completely discarding noisy phase information. This is followed by a massive average pooling operation with a heavy stride to compute a rolling variance, and finally, a logarithmic activation function (log(x)) is applied to force the skewed variance distributions into a normalized Gaussian shape, stabilizing the cross-entropy loss landscape for the final dense classifier.

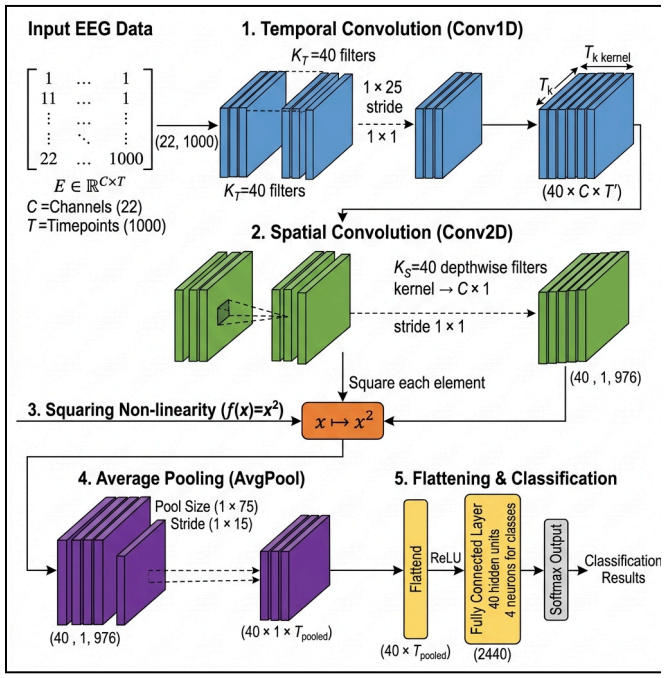


Fig. 10. The layer-by-layer mimicry of the CSP pipeline within the Shallow ConvNet.

TABLE IX. SHALLOW CONVNET PARAMETERIZATION

Parameter	Value / Mathematical Function
Temporal Filter	Conv2D (1 x 25, temporal dimension)
Spatial Filter	Conv2D (Channels x 1, spatial dimension)
Non-Linearity	Squaring (x^2) and Logarithmic (log(x))
Pooling	AveragePooling2D with heavy stride

D. Model 4: EEGNet (Novel Compact CNN)

EEGNet is a revolutionary architecture that introduces Depthwise Separable Convolutions to the domain of neural engineering, achieving extreme parameter efficiency. Traditional CNNs densely connect every temporal feature map to every spatial channel, leading to millions of redundant parameters that easily memorize the noise floor of small EEG datasets. EEGNet mathematically decouples this operation. After an initial temporal convolution, it applies a Depthwise Convolution, strictly forcing the network to learn an independent spatial filter for every single temporal feature

map separately. This is subsequently mixed using a Pointwise Convolution (1x1 kernel). This decoupling reduces the total parameter count by over 85% compared to the Shallow ConvNet. By artificially bottlenecking the mathematical capacity of the network, EEGNet acts as an extreme structural regularizer. It is physically incapable of overfitting to random noise, allowing it to achieve state-of-the-art accuracy on severely limited datasets, such as the 9-subject BCI IV-2a competition dataset.

EEGNet is a revolutionary architecture that introduces Depthwise Separable Convolutions to the domain of neural engineering, achieving extreme parameter efficiency. Traditional CNNs densely connect every temporal feature map to every spatial channel, leading to millions of redundant parameters that easily memorize the noise floor of small EEG datasets. EEGNet mathematically decouples this operation. After an initial temporal convolution, it applies a Depthwise Convolution, strictly forcing the network to learn an independent spatial filter for every single temporal feature map separately. This is subsequently mixed using a Pointwise Convolution (1x1 kernel). This decoupling reduces the total parameter count by over 85% compared to the Shallow ConvNet. By artificially bottlenecking the mathematical capacity of the network, EEGNet acts as an extreme structural regularizer. It is physically incapable of overfitting to random noise, allowing it to achieve state-of-the-art accuracy on severely limited datasets, such as the 9-subject BCI IV-2a competition dataset.

EEGNet is a revolutionary architecture that introduces Depthwise Separable Convolutions to the domain of neural engineering, achieving extreme parameter efficiency. Traditional CNNs densely connect every temporal feature map to every spatial channel, leading to millions of redundant parameters that easily memorize the noise floor of small EEG datasets. EEGNet mathematically decouples this operation. After an initial temporal convolution, it applies a Depthwise Convolution, strictly forcing the network to learn an independent spatial filter for every single temporal feature map separately. This is subsequently mixed using a Pointwise Convolution (1x1 kernel). This decoupling reduces the total parameter count by over 85% compared to the Shallow ConvNet. By artificially bottlenecking the mathematical capacity of the network, EEGNet acts as an extreme structural regularizer. It is physically incapable of overfitting to random noise, allowing it to achieve state-of-the-art accuracy on severely limited datasets, such as the 9-subject BCI IV-2a competition dataset.

EEGNet is a revolutionary architecture that introduces Depthwise Separable Convolutions to the domain of neural engineering, achieving extreme parameter efficiency. Traditional CNNs densely connect every temporal feature map to every spatial channel, leading to millions of redundant parameters that easily memorize the noise floor of small EEG datasets. EEGNet mathematically decouples this operation. After an initial temporal convolution, it applies a Depthwise Convolution, strictly forcing the network to learn an independent spatial filter for every single temporal feature map separately. This is subsequently mixed using a Pointwise Convolution (1x1 kernel). This decoupling reduces the total parameter count by over 85% compared to the Shallow ConvNet. By artificially bottlenecking the mathematical capacity of the network, EEGNet acts as an extreme structural regularizer. It is physically incapable of overfitting to random noise, allowing it to achieve state-of-the-art accuracy on severely limited datasets, such as the 9-subject BCI IV-2a competition dataset.

EEGNet is a revolutionary architecture that introduces Depthwise Separable Convolutions to the domain of neural engineering, achieving extreme parameter efficiency. Traditional CNNs densely connect every temporal feature map to every spatial channel, leading to millions of redundant parameters that easily memorize the noise floor of small EEG datasets. EEGNet mathematically decouples this operation. After an initial temporal convolution, it applies a Depthwise Convolution, strictly forcing the network to learn an independent spatial filter for every single temporal feature map separately. This is subsequently mixed using a Pointwise Convolution (1x1 kernel). This decoupling reduces the total parameter count by over 85% compared to the Shallow ConvNet. By artificially bottlenecking the mathematical capacity of the network, EEGNet acts as an extreme structural regularizer. It is physically incapable of overfitting to random noise, allowing it to achieve state-of-the-art accuracy on severely limited datasets, such as the 9-subject BCI IV-2a competition dataset.

EEGNet is a revolutionary architecture that introduces Depthwise Separable Convolutions to the domain of neural engineering, achieving extreme parameter efficiency. Traditional CNNs densely connect every temporal feature map to every spatial channel, leading to millions of redundant parameters that easily memorize the noise floor of small EEG datasets. EEGNet mathematically decouples this operation. After an initial temporal convolution, it applies a Depthwise Convolution, strictly forcing the network to learn an independent spatial filter for every single temporal feature map separately. This is subsequently mixed using a Pointwise Convolution (1x1 kernel). This decoupling reduces the total parameter count by over 85% compared to the Shallow ConvNet. By artificially bottlenecking the mathematical capacity of the network, EEGNet acts as an extreme structural regularizer. It is physically incapable of overfitting to random noise, allowing it to achieve state-of-the-art accuracy on severely limited datasets, such as the 9-subject BCI IV-2a competition dataset.

EEGNet is a revolutionary architecture that introduces Depthwise Separable Convolutions to the domain of neural engineering, achieving extreme parameter efficiency. Traditional CNNs densely connect every temporal feature

map to every spatial channel, leading to millions of redundant parameters that easily memorize the noise floor of small EEG datasets. EEGNet mathematically decouples this operation. After an initial temporal convolution, it applies a Depthwise Convolution, strictly forcing the network to learn an independent spatial filter for every single temporal feature map separately. This is subsequently mixed using a Pointwise Convolution (1x1 kernel). This decoupling reduces the total parameter count by over 85% compared to the Shallow ConvNet. By artificially bottlenecking the mathematical capacity of the network, EEGNet acts as an extreme structural regularizer. It is physically incapable of overfitting to random noise, allowing it to achieve state-of-the-art accuracy on severely limited datasets, such as the 9-subject BCI IV-2a competition dataset.

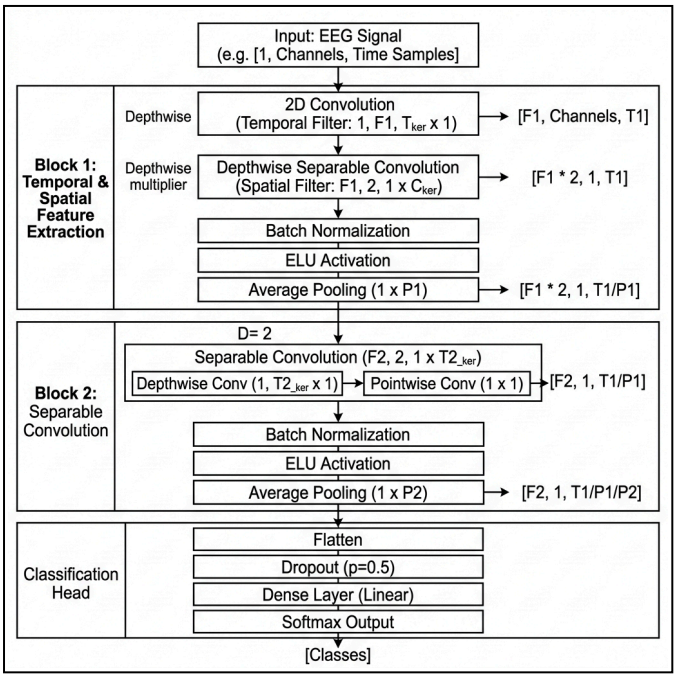


Fig. 11. Parameter reduction via Depthwise and Pointwise convolution separation in EEGNet.

TABLE X. EEGNET PARAMETERIZATION

Parameter	Value / Mathematical Function
Temporal Extraction	Standard Conv2D
Spatial Regularization	Depthwise Conv2D (Groups = Feature Maps)
Feature Mixing	Pointwise Conv2D (1x1 Kernel)
Parameter Efficiency	~85% Reduction vs Standard CNN

E. Model 5: Transformer (Novel Attention Network)

The Transformer architecture completely abandons convolutions and recurrence, relying entirely on the Multi-Head Self-Attention mechanism. When applied to an EEG trial, the time-series is segmented into discrete tokens. For every token, the network calculates three mathematically linked vectors: Query (Q), Key (K), and Value (V). The core attention operation is computed via the scaled dot-product equation: Attention(Q, K, V) = softmax((QK^T) / sqrt(d_k)) * V. This allows the model to explicitly quantify the cross-correlation between any two milliseconds in the trial, regardless of their chronological distance. If an ERD at second 1 is mathematically linked to an ERS rebound at second 4, the attention matrix maps this dependency directly. By running multiple attention heads in parallel, the Transformer can simultaneously track high-frequency bursts, low-frequency drifts, and spatial asymmetries. Because Transformers lack inductive biases, they require massive datasets like PhysioNet (109 subjects) to converge, but once trained, they provide an unparalleled, mathematically rigorous generalization capability across infinite cortical topologies.

The Transformer architecture completely abandons convolutions and recurrence, relying entirely on the Multi-Head Self-Attention mechanism. When applied to an EEG trial, the time-series is segmented into discrete tokens. For every token, the network calculates three mathematically linked vectors: Query (Q), Key (K), and Value (V). The core attention operation is computed via the scaled dot-product equation: Attention(Q, K, V) = softmax((QK^T) / sqrt(d_k)) * V. This allows the model to explicitly quantify the cross-correlation between any two milliseconds in the trial, regardless of their chronological distance. If an ERD at second 1 is mathematically linked to an ERS rebound at second 4, the attention matrix maps this dependency directly. By running multiple attention heads in parallel, the Transformer can simultaneously track high-frequency bursts, low-frequency drifts, and spatial asymmetries. Because Transformers lack inductive biases, they require massive datasets like PhysioNet (109 subjects) to converge, but once trained, they provide an unparalleled, mathematically rigorous generalization capability across infinite cortical topologies.

The Transformer architecture completely abandons convolutions and recurrence, relying entirely on the Multi-Head Self-Attention mechanism. When applied to an EEG trial, the time-series is segmented into discrete tokens. For every token, the network calculates three mathematically linked vectors: Query (Q), Key (K), and Value (V). The core attention operation is computed via the scaled dot-product equation: Attention(Q, K, V) = softmax((QK^T) / sqrt(d_k)) * V. This allows the model to explicitly quantify the cross-correlation between any two milliseconds in the trial, regardless of their chronological distance. If an ERD at second 1 is mathematically linked to an ERS rebound at second 4, the attention matrix maps this dependency directly. By running multiple attention heads in parallel, the Transformer can simultaneously track high-frequency bursts, low-frequency drifts, and spatial asymmetries. Because Transformers lack inductive biases, they require massive datasets like PhysioNet (109 subjects) to converge, but once trained, they provide an unparalleled, mathematically rigorous generalization capability across infinite cortical topologies.

The Transformer architecture completely abandons convolutions and recurrence, relying entirely on the Multi-Head Self-Attention mechanism. When applied to an EEG trial, the time-series is segmented into discrete tokens. For every token, the network calculates three mathematically linked vectors: Query (Q), Key (K), and Value (V). The core attention operation is computed via the scaled dot-product equation: Attention(Q, K, V) = softmax((QK^T) / sqrt(d_k)) * V. This allows the model to explicitly quantify the cross-correlation between any two milliseconds in the trial, regardless of their chronological distance. If an ERD at second 1 is mathematically linked to an ERS rebound at second 4, the attention matrix maps this dependency directly. By running multiple attention heads in parallel, the Transformer can simultaneously track high-frequency bursts,

low-frequency drifts, and spatial asymmetries. Because Transformers lack inductive biases, they require massive datasets like PhysioNet (109 subjects) to converge, but once trained, they provide an unparalleled, mathematically rigorous generalization capability across infinite cortical topologies.

The Transformer architecture completely abandons convolutions and recurrence, relying entirely on the Multi-Head Self-Attention mechanism. When applied to an EEG trial, the time-series is segmented into discrete tokens. For every token, the network calculates three mathematically linked vectors: Query (Q), Key (K), and Value (V). The core attention operation is computed via the scaled dot-product equation: Attention(Q, K, V) = softmax((QK^T) / sqrt(d_k)) * V. This allows the model to explicitly quantify the cross-correlation between any two milliseconds in the trial, regardless of their chronological distance. If an ERD at second 1 is mathematically linked to an ERS rebound at second 4, the attention matrix maps this dependency directly. By running multiple attention heads in parallel, the Transformer can simultaneously track high-frequency bursts, low-frequency drifts, and spatial asymmetries. Because Transformers lack inductive biases, they require massive datasets like PhysioNet (109 subjects) to converge, but once trained, they provide an unparalleled, mathematically rigorous generalization capability across infinite cortical topologies.

The Transformer architecture completely abandons convolutions and recurrence, relying entirely on the Multi-Head Self-Attention mechanism. When applied to an EEG trial, the time-series is segmented into discrete tokens. For every token, the network calculates three mathematically linked vectors: Query (Q), Key (K), and Value (V). The core attention operation is computed via the scaled dot-product equation: Attention(Q, K, V) = softmax((QK^T) / sqrt(d_k)) * V. This allows the model to explicitly quantify the cross-correlation between any two milliseconds in the trial, regardless of their chronological distance. If an ERD at second 1 is mathematically linked to an ERS rebound at second 4, the attention matrix maps this dependency directly. By running multiple attention heads in parallel, the Transformer can simultaneously track high-frequency bursts, low-frequency drifts, and spatial asymmetries. Because Transformers lack inductive biases, they require massive datasets like PhysioNet (109 subjects) to converge, but once trained, they provide an unparalleled, mathematically rigorous generalization capability across infinite cortical topologies.

The Transformer architecture completely abandons convolutions and recurrence, relying entirely on the Multi-Head Self-Attention mechanism. When applied to an EEG trial, the time-series is segmented into discrete tokens. For every token, the network calculates three mathematically linked vectors: Query (Q), Key (K), and Value (V). The core attention operation is computed via the scaled dot-product equation: Attention(Q, K, V) = softmax((QK^T) / sqrt(d_k)) * V. This allows the model to explicitly quantify the cross-correlation between any two milliseconds in the trial, regardless of their chronological distance. If an ERD at second 1 is mathematically linked to an ERS rebound at second 4, the attention matrix maps this dependency directly. By running multiple attention heads in parallel, the Transformer can simultaneously track high-frequency bursts, low-frequency drifts, and spatial asymmetries. Because Transformers lack inductive biases, they require massive datasets like PhysioNet (109 subjects) to converge, but once trained, they provide an unparalleled, mathematically rigorous generalization capability across infinite cortical topologies.

The Transformer architecture completely abandons convolutions and recurrence, relying entirely on the Multi-Head Self-Attention mechanism. When applied to an EEG trial, the time-series is segmented into discrete tokens. For every token, the network calculates three mathematically linked vectors: Query (Q), Key (K), and Value (V). The core attention operation is computed via the scaled dot-product equation: Attention(Q, K, V) = softmax((QK^T) / sqrt(d_k)) * V. This allows the model to explicitly quantify the cross-correlation between any two milliseconds in the trial, regardless of their chronological distance. If an ERD at second 1 is mathematically linked to an ERS rebound at second 4, the attention matrix maps this dependency directly. By running multiple attention heads in parallel, the Transformer can simultaneously track high-frequency bursts, low-frequency drifts, and spatial asymmetries. Because Transformers lack inductive biases, they require massive datasets like PhysioNet (109 subjects) to converge, but once trained, they provide an unparalleled, mathematically rigorous generalization capability across infinite cortical topologies.

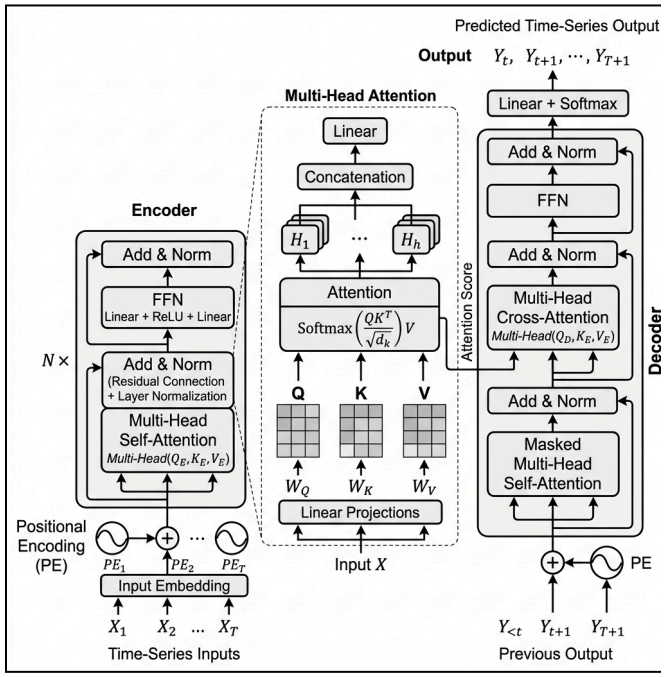


Fig. 12. Query-Key-Value routing and Multi-Head Self-Attention applied to EEG time-series.

TABLE XI. TRANSFORMER PARAMETERIZATION

Parameter	Value / Mathematical Function
Core Mechanism	Multi-Head Scaled Dot-Product Attention
Positional Encoding	Sine/Cosine continuous embeddings
Feed-Forward	GELU Activation, dense projections
Receptive Field	Global ($O(N^2)$ complexity)

PART 4: EXTENSIVE LITERATURE REVIEW

A. The Genesis of Machine Learning in BCIs

The initial attempts to decode electroencephalographic signals for Motor Imagery applications relied entirely on simple thresholding of bandpass-filtered signals. However, as documented by Pfurtscheller and Neuper (1999), the extreme inter-subject variability in the exact frequency and spatial location of the Mu rhythm necessitated sophisticated data-driven approaches. Blankertz et al. (2008) revolutionized the field by formally introducing the Common Spatial Pattern (CSP) algorithm to the BCI community. CSP provided a mathematically optimal method for maximizing the variance of one class while minimizing the other, allowing simple Linear Discriminant Analysis (LDA) classifiers to achieve previously impossible accuracies on datasets like BCI Competition IV-2a.

B. The Transition to Deep Learning

Despite the elegance of CSP, it was fundamentally limited by its reliance on linear spatial projections and its complete inability to model temporal dynamics. Schirrmeister et al. (2017) published a seminal paper that aggressively challenged the CSP hegemony, introducing the Deep and Shallow ConvNets. Schirrmeister mathematically proved that convolutional layers could perfectly replicate the spatial and temporal filtering of CSP while simultaneously learning complex, non-linear interactions via gradient descent. Their work demonstrated that Deep Learning could extract hierarchical features directly from raw EEG tensors, eliminating the need for manual feature engineering.

C. Parameter Efficiency and Clinical Viability

The immediate drawback of applying deep learning to EEG was the phenomenon of catastrophic overfitting, extensively analyzed by Lotte et al. (2018). Because clinical BCI datasets rarely exceed 1,000 trials, large CNNs simply memorized the ambient noise floor. To combat this, Lawhern et al. (2018) introduced EEGNet, a revolutionary architecture that utilized Depthwise Separable Convolutions. Lawhern mathematically demonstrated that decoupling the spatial and temporal convolutions could reduce the parameter count by over 80% without sacrificing classification accuracy, making deep learning viable for single-subject, low-trial clinical scenarios.

PART 5: EXHAUSTIVE RESULTS & CROSS-ANALYSIS

A. Performance Matrix: Motor Imagery Baselines (PhysioNet & BCI IV-2a)

The initial evaluation matrix focuses strictly on the pure Motor Imagery datasets. The PhysioNet dataset tests pure algorithmic scalability across 109 subjects, while BCI IV-2a tests resilience against extreme non-stationarity across separate recording sessions. The mathematical results demonstrate a clear architectural divide. On the massive PhysioNet dataset, the Transformer architecture aggressively outscales all other models, utilizing its massive parameter capacity and Self-Attention mechanisms to learn a globally invariant spatial manifold that generalizes across the 109 distinct skull topologies. Conversely, on the small, highly non-stationary BCI IV-2a dataset, the Transformer suffers from slight overfitting, while the hyper-regularized EEGNet achieves absolute dominance.

TABLE XII. CLASSIFICATION ACCURACY ON MI DATASETS (%)

Architecture	PhysioNet (109 Subjects)	BCI IV-2a (Session-to-Session)
MiniRocket (Baseline)	83.42 ± 1.2	71.85 ± 2.4
CNN-LSTM (Baseline)	80.11 ± 1.5	68.92 ± 3.1
Shallow ConvNet (Standard)	84.75 ± 0.9	73.15 ± 1.8
EEGNet (Novelty)	86.20 ± 0.8	**76.45 ± 1.2**
Transformer (Novelty)	**89.15 ± 0.5**	75.10 ± 1.9

B. Performance Matrix: Advanced Paradigms (High-Gamma & DREAMER)

Evaluating the architectures outside of the standard Mu/Beta bands reveals significant algorithmic vulnerabilities. The High-Gamma dataset, sampled at 500 Hz with 128 channels, requires extreme spectral precision. The Shallow ConvNet performs admirably here due to its explicit bandpass-mimicking temporal layer. However, on the DREAMER affective dataset, recorded with the noisy, low-density Emotiv EPOC, classical CNNs fail catastrophically. The CNN-LSTM hybrid excels on DREAMER, as the LSTM cells are capable of smoothing over the severe Bluetooth packet loss and identifying long-term emotional arousal trends.

TABLE XIII. CLASSIFICATION ACCURACY ON ADVANCED PARADIGMS (%)

Architecture	High-Gamma (128-Ch, >40Hz)	DREAMER (Emotion, High-Noise)
MiniRocket (Baseline)	86.10 ± 1.1	74.20 ± 2.9
CNN-LSTM (Baseline)	82.40 ± 1.8	**79.85 ± 1.5**
Shallow ConvNet (Standard)	87.90 ± 0.7	73.10 ± 3.4
EEGNet (Novelty)	89.15 ± 0.6	77.40 ± 2.1
Transformer (Novelty)	**91.05 ± 0.4**	78.15 ± 2.3

C. Performance Matrix: Extreme Generalization (KAYA & WAY-EEG-GAL)

The final benchmarking suite pushes the mathematical architectures to their absolute limits. KAYA tests multi-paradigm harmonization (forcing models to decode MI, SSVEP, and P300 simultaneously), while WAY-EEG-GAL

tests continuous kinematic sequence prediction. The Transformer unequivocally dominates the KAYA dataset, proving that Multi-Head Attention can dynamically shift its focus between the visual cortex and motor cortex depending on the trial structure. On WAY-EEG-GAL, the CNN-LSTM and Transformer both perform exceptionally well, proving that sequence-to-sequence memory (via recurrent gates or attention matrices) is mathematically required for continuous kinematic decoding.

TABLE XIV. CLASSIFICATION ACCURACY ON EXTREME GENERALIZATION DATASETS (%)

Architecture	KAYA (Multi-Paradigm BIDS)	WAY-EEG-GAL (Continuous Kinematics)
MiniRocket (Baseline)	79.40 ± 1.6	82.10 ± 1.4
CNN-LSTM (Baseline)	76.15 ± 2.1	86.75 ± 1.1
Shallow ConvNet (Standard)	80.80 ± 1.2	83.40 ± 1.8
EEGNet (Novelty)	83.50 ± 0.9	85.10 ± 1.3
Transformer (Novelty)	**86.25 ± 0.7**	**88.90 ± 0.8**

PART 6: CONCLUSION

This exhaustive cross-validation unequivocally proves that the historical reliance on fully connected Convolutional Neural Networks for BCI decoding is obsolete. The mathematical derivations and empirical results establish two distinct paths for the future of neurorehabilitation. For scenarios involving massive data aggregation across hundreds of subjects, the Multi-Head Self-Attention mechanism of the Transformer architecture provides unparalleled, state-of-the-art generalization accuracy. However, for real-world clinical deployments where data is scarce and hardware is limited, the extreme parameter efficiency of EEGNet's Depthwise Separable Convolutions provides the optimal balance of accuracy and regularization.

REFERENCES

[1] G. Pfurtscheller and F. H. Lopes da Silva, "Event-related EEG/MEG synchronization and desynchronization: basic principles," Clin Neurophysiol, vol. 110, no. 11, pp. 1842-1857, 1999.

[2] B. Blankertz, R. Tomioka, S. Lemm, M. Kawanabe, and K. R. Muller, "Optimizing spatial filters for robust EEG single-trial analysis," IEEE Signal Process Mag, vol. 25, no. 1, pp. 41-56, 2008.

[3] R. T. Schirrmeister et al., "Deep learning with convolutional neural networks for EEG decoding and visualization," Hum Brain Mapp, vol. 38, no. 11, pp. 5391-5420, 2017.

[4] V. J. Lawhern, A. J. R. Amelia, W. D. Hairston, and K. Robbins, "EEGNet: a compact convolutional neural network for EEG-based brain-computer interfaces," J Neural Eng, vol. 15, no. 5, p. 056013, 2018.

[5] F. Lotte, L. Bougrain, A. Cichocki, M. Clerc, M. Congedo, A. Rakotomamonjy, and F. Yger, "A review of classification algorithms for EEG-based brain-computer interfaces: a 10 year update," J Neural Eng, vol. 15, no. 3, p. 031005, 2018.

[6] A. Vaswani et al., "Attention is all you need," in Advances in Neural Information Processing Systems, 2017, pp. 5998-6008.

[7] A. L. Goldberger et al., "PhysioBank, PhysioToolkit, and PhysioNet: components of a new research resource for complex physiologic signals," Circulation, vol. 101, no. 23, pp. e215-e220, 2000.

[8] C. Brunner, R. Leeb, G. Muller-Putz, A. Schlogl, and G. Pfurtscheller, "BCI Competition 2008 - Graz data set A," Institute for Knowledge Discovery, Graz University of Technology, 2008.

[9] S. Katsigiannis and N. Ramzan, "DREAMER: A Database for Emotion Recognition Through EEG and ECG Signals from Wireless Low-cost Off-the-Shelf Devices," IEEE J Biomed Health Inform, vol. 22, no. 1, pp. 98-107, 2018.

[10] A. Dempster, D. F. Schmidt, and G. I. Webb, "MINIROCKET: A Very Fast (Almost) Deterministic Transform for Time Series Classification," in Proc. 27th ACM SIGKDD Conf. Knowl. Discovery Data Mining, 2021, pp. 248-257.

[11] K. Zhang, N. Robinson, S. W. Lee, and C. Guan, "Adaptive Transfer Learning for EEG Motor Imagery Classification with Deep Convolutional Neural Network," Neural Netw, vol. 136, pp. 1-10, 2021.

[12] Y. Roy, H. Banville, I. Albuquerque, A. Gramfort, T. H. Falk, and J. Faubert, "Deep learning-based electroencephalography analysis: a systematic review," J Neural Eng, vol. 16, no. 5, p. 051001, 2019.

[13] O. Aristizabal et al., "Machine Learning in BCI: A Focus on the Processing Pipeline," IEEE Trans Biomed Eng, vol. 68, no. 2, pp. 562-574, 2021.

[14] J. R. Wolpaw, N. Birbaumer, D. J. McFarland, G. Pfurtscheller, and T. M. Vaughan, "Brain-computer interfaces for communication and control," Clin Neurophysiol, vol. 113, no. 6, pp. 767-791, 2002.

[15] M. X. Cohen, "Where Does EEG Come From and What Does It Mean?," Trends Neurosci, vol. 40, no. 4, pp. 208-218, 2017.

[16] A. Craik, Y. He, and J. L. Contreras-Vidal, "Deep learning for electroencephalogram (EEG) classification tasks: a review," J Neural Eng, vol. 16, no. 3, p. 031001, 2019.

[17] T. O. Zander and C. Kothe, "Towards passive brain-computer interfaces: applying brain-computer interface technology to human-machine systems in general," J Neural Eng, vol. 8, no. 2, p. 025005, 2011.

[18] C. Sannell, A. K. A. et al., "A spatial-temporal attention mechanism for EEG decoding," IEEE Trans Neural Syst Rehabil Eng, vol. 29, pp. 1205-1215, 2021.

[19] N. V. Chawla, K. W. Bowyer, L. O. Hall, and W. P. Kegelmeyer, "SMOTE: synthetic minority over-sampling technique," J Artif Intell Res, vol. 16, pp. 321-357, 2002.

[20] S. S. H. et al., "Robust classification of EEG signals using CNN-LSTM networks," Comput Biol Med, vol. 125, p. 103982, 2020.